

Operating Systems

Associate Prof. Yongkun Li

中科大-计算机学院 副教授

<http://staff.ustc.edu.cn/~ykli>


Ch5

Process Communication & Synchronization

-Part 3

Story so far...

- For shared memory and files, concurrent access may yield unpredictable outcomes
 - Race condition
- To avoid race condition, **mutual exclusion** must be guaranteed
 - Critical section
 - Implementations (entry/exit)
 - Hardware instructions
 - Disabling interrupts
 - Strict alternation
 - Peterson's solution
 - Mutex lock
 - **Semaphore**



Semaphore Usage

- Semaphore can be used for
 - Mutual exclusion (binary semaphore)
 - Process synchronization (counting semaphore may be needed)
- How to do **process synchronization** w/ semaphore?
 - Mutual exclusion + coordination (multiple semaphores)
 - Careless design may lead to other issues
 - Deadlock

Topics

Deadlock

- ☐ Concept
- ☐ Necessary conditions
- ☐ Characterization
- ☐ Solutions

Classic problems


- ☐ Producer-consumer problem
- ☐ Dining philosopher problem
- ☐ Reader-writer problem

The Deadlock Problem

Classic IPC problems

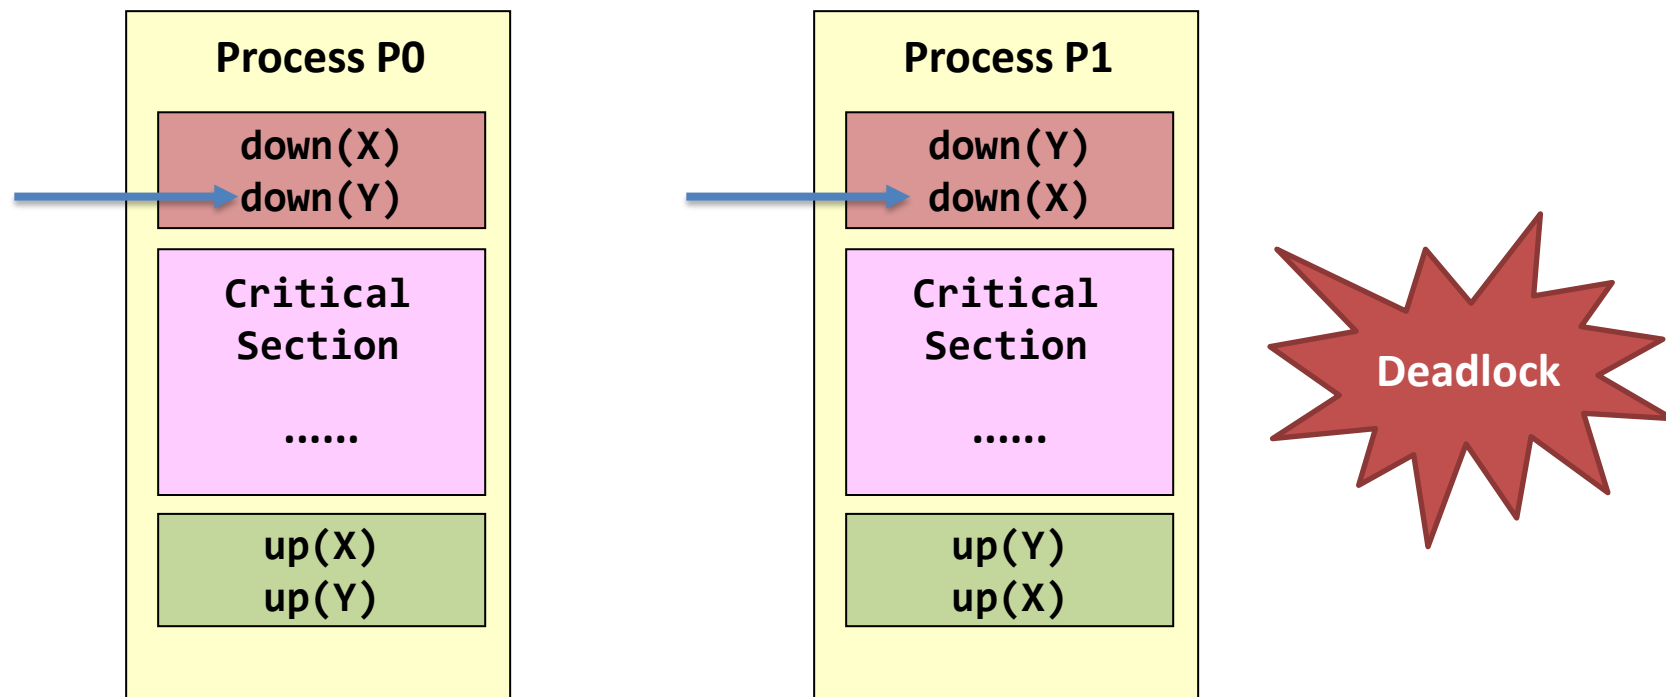
- Producer-consumer problem
- Dining philosopher problem
- Reader-writer problem

Let's teach them
not to fight.



Deadlock Example

- Problems when using semaphore



Scenario: P0 must wait until P1 executes `up(Y)`, P1 must wait until P0 executes `up(X)`

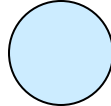
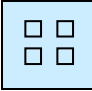
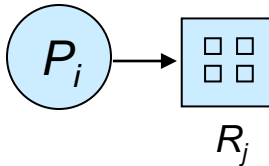
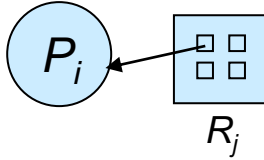
Deadlock Requirements

- **Requirement #1: Mutual Exclusion.**
 - Only one process at a time can use a resource
- **Requirement #2. Hold and wait.**
 - A process must be holding at least one resource and waiting to acquire additional resources held by other processes



Deadlock Requirements

- **Requirement #3: No preemption.**
 - A resource can be released only voluntarily by the process holding it after that process has completed its task
- **Requirement #4. Circular wait.**
 - There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 waits for P_1 , P_1 waits for P_2 , ..., P_{n-1} waits for P_n , P_n waits for P_0

How to Handle Deadlocks


- Deadlock characterization: Deadlocks can be described using **resource-allocation graph**
 - Set V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$: processes 
 - $R = \{R_1, R_2, \dots, R_m\}$: all resource types (each type may have multiple instances) 
 - Set E
 - **request edge** – directed edge $P_i \rightarrow R_j$ 
 - **assignment edge** – directed edge $R_j \rightarrow P_i$ 

Examples



How to Handle Deadlocks

- **Detect** deadlock and recover
 - Case 1: Each resource has one instance
 - Resource-allocation graph: detect the existence of a cycle



No cycles


No deadlock

Contains a cycle


Case 1: only one instance per resource

Case 2: several instances per resource type: possible deadlock

Examples



Deadlock



No deadlock

How to Handle Deadlocks

- **Detect** deadlock and recover

- Case 2: Each resource has multiple instances

- Matrix method: four data structures

- Existing (total) resources (m types): (E_1, E_2, \dots, E_m)

- Available resources: (A_1, A_2, \dots, A_m)

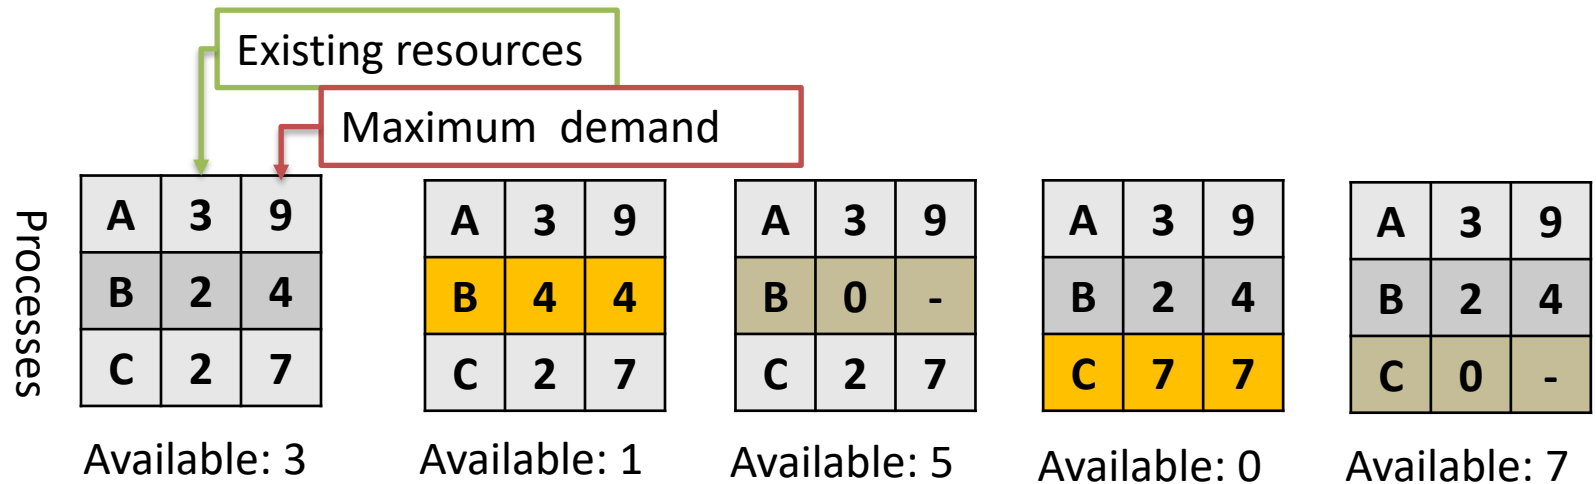
- Allocation matrix: $\begin{bmatrix} C_{11} & \cdots & C_{1m} \\ \vdots & \ddots & \vdots \\ C_{n1} & \cdots & C_{nm} \end{bmatrix}$ (C_{ij} : # of type- j resources held by process i)

- Request matrix: $\begin{bmatrix} R_{11} & \cdots & R_{1m} \\ \vdots & \ddots & \vdots \\ R_{n1} & \cdots & R_{nm} \end{bmatrix}$ (R_{ij} : # of type- j resources requested by process i)

- Repeatedly check P_i s.t. $R_i \leq A$? (P_i can be satisfied)
 - ✓ Yes: $A = A + C_i$ (release resources)
 - ✓ No: End (remaining processes are deadlocked)

How to Handle Deadlocks

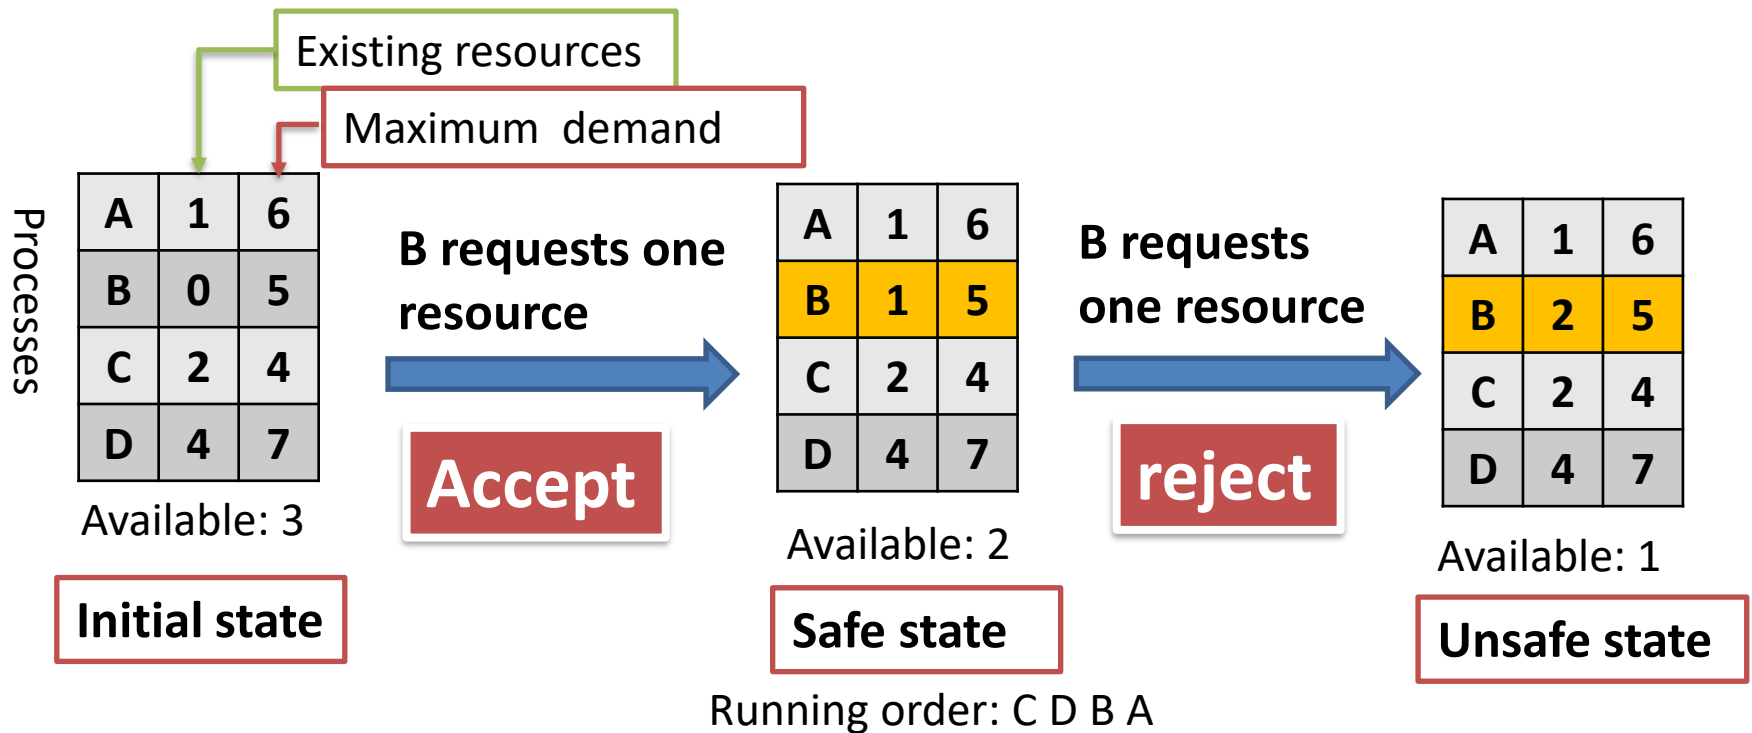
- **Prevent/avoid** deadlocks: Banker's algorithm
 - Idea: check system state defined by (E, A, C, R)
 - **Safe state**: exist one running sequence to guarantee that all processes' demand can be satisfied



- **Unsafe state**: Not exist any sequence to guarantee the demand
 - It is not deadlock (it can still run for some time/processes may release some resources)

How to Handle Deadlocks

- **Prevent/avoid** deadlocks: Banker's algorithm
 - For each request: safe (accept), unsafe (reject)



The algorithm can also be extended to the case of multiple resources, but it needs to know the demand

How to Handle Deadlocks


- **Ignore** the problem and pretend that deadlocks never occur (stop functioning and restart manually)
 - 鸵鸟算法（假装没发生）
 - Used by most operating systems, including UNIX and windows
 - Deadlocks occur infrequently, avoiding/detecting it is expensive
- A deadlock-free solution does not eliminate **starvation**

The Deadlock Problem

Classic IPC problems

- Producer-consumer problem
- Dining philosopher problem
- Reader-writer problem

Let's teach them
not to fight.



What are the problems?

- All the IPC classical problems use **semaphores** to fulfill the synchronization requirements.


| | Properties | Examples |
|---------------------------|---|-----------------------------|
| Producer-Consumer Problem | Two classes of processes: <u>producer</u> and <u>consumer</u> ; At least one producer and one consumer. | FIFO buffer, such as pipe. |
| Dining Philosophy Problem | They are all running the same program; At least two processes. | Cross-road traffic control. |
| Reader-Writer Problem | Two classes of processes: <u>reader</u> and <u>writer</u> . No limit on the number of the processes of each class. | Database. |

The Deadlock Problem

Classic IPC problems

- **Producer-consumer problem**
- Dining philosopher problem
- Reader-writer problem

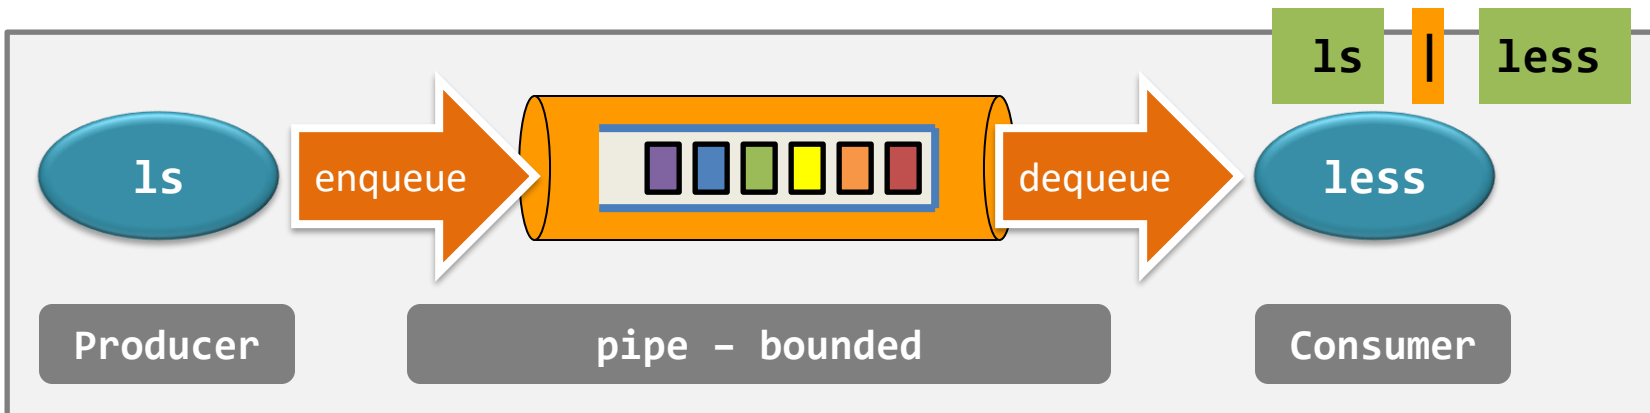
Let's teach them
not to fight.



Producer-consumer problem – recall

- Also known as the **bounded-buffer problem**.

| | |
|---------------------------|--|
| A bounded buffer | <ul style="list-style-type: none">-It is a shared object;-Its size is bounded, say N slots.-It is a queue (imagine that it is an array implementation of queue). |
| A producer process | <ul style="list-style-type: none">-It produces a unit of data, and-writes that a piece of data to the tail of the buffer at one time. |
| A consumer process | <ul style="list-style-type: none">-It removes a unit of data from the head of the bounded buffer at one time. |



Producer-consumer problem – recall

Producer-consumer requirement #1

When the producer wants to
(a) put a new item in the buffer, but
(b) **the buffer is already full...**

Then,

- (1) **The producer should be suspended**, and
- (2) **The consumer should wake the producer up** after she has dequeued an item.

Producer-consumer requirement #2

When the consumer wants to
(a) consumes an item from the buffer, but
(b) **the buffer is empty...**

Then,

- (1) **The consumer should be suspended**, and
- (2) **The producer should wake the consumer up** after she has enqueued an item.

Producer-consumer problem

- Pipe is working fine. Is it enough?
 - What if we cannot use pipes?
 - Say, there are 2 producers and 2 consumers without any parent-child relationships?
 - Then, **the kernel can't protect you with a pipe.**
- In the following, we revisit the producer-consumer problem with the use of shared objects and semaphores, instead of pipe.

Design – Semaphores

- **ISSUE #1: Mutual Exclusion.**

Solution: one binary semaphore (mutex)

- **ISSUE #2: Synchronization (coordination).**

- Remember the two requirements:
 - Insert an item when it is not FULL
 - Consume an item when it is not EMPTY
- Can we use a binary semaphore?

Solution: two counting semaphores (full & empty)

Producer-consumer problem – solution

Note

The functions “`insert_item()`” and “`remove_item()`” are accessing the bounded buffer (codes in critical section).

The size of the bounded buffer is “N”.

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6
7
8         insert_item(item);
9
10
11     }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5
6
7         item = remove_item();
8
9
10        consume_item(item);
11    }
12 }
```


Producer-consumer problem – solution

Note

Mutual exclusion requirement

Synchronization requirement

Shared object

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6
7         insert_item(item);
8
9
10
11     }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5
6
7         item = remove_item();
8
9
10        consume_item(item);
11    }
12 }
```

Producer-consumer problem – Understanding

Why we need three semaphores, “empty”, “full”, “mutex”?

Shared object

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6
7
8         insert_item(item);
9
10
11     }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5
6
7         item = remove_item();
8
9
10        consume_item(item);
11    }
12 }
```

Producer-consumer problem – Understanding

Why we need three semaphores, “empty”, “full”, “mutex”?

mutex:

What is its purpose?

Why is the initial value of mutex 1?

Shared object

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6
7         down(&mutex);
8         insert_item(item);
9         up(&mutex);
10    }
11 }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5
6         down(&mutex);
7         item = remove_item();
8         up(&mutex);
9
10        consume_item(item);
11    }
12 }
```

Producer-consumer problem – Understanding

Why we need three semaphores, “empty”, “full”, “mutex”?

mutex:

what is its purpose?

Why is the initial value of mutex 1?

Shared object

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6
7         down(&mutex);
8         insert_item(item);
9         up(&mutex);
10    }
11 }
12 }
```

The “**mutex**” stands for mutual exclusion.

- **down()** and **up()** statements are the entry and the exit of the critical section, respectively.

What is the meaning of the initial value 1?

Producer-consumer problem – Understanding

Why we need three semaphores, “empty”, “full”, “mutex”?

How about “full” and “empty”?

Shared object

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6         down(&empty);
7         down(&mutex);
8         insert_item(item);
9         up(&mutex);
10        up(&full);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         down(&full);
6         down(&mutex);
7         item = remove_item();
8         up(&mutex);
9         up(&empty);
10        consume_item(item);
11    }
12 }
```

Producer-consumer problem – Understanding

- The two variables are not for mutual exclusion, but for **process synchronization**.
 - “*Process synchronization*” means **to coordinate** the set of processes so as to produce meaningful output.

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6         down(&empty);
7         down(&mutex);
8         insert_item(item);
9         up(&mutex);
10        up(&full);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         down(&full);
6         down(&mutex);
7         item = remove_item();
8         up(&mutex);
9         up(&empty);
10        consume_item(item);
11    }
12 }
```

Producer-consumer problem – Understanding

For “empty”,

- Its initial value is N;
- It decrements by 1 in each iteration.
- When it reaches 0, the producers sleeps.

So, does it sound like one of the requirements?

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

The consumer wakes the producer up when it finds “empty” is 0.

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6         down(&empty);
7         down(&mutex);
8         insert_item(item);
9         up(&mutex);
10        up(&full);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         down(&full);
6         down(&mutex);
7         item = remove_item();
8         up(&mutex);
9         up(&empty);
10        consume_item(item);
11    }
12 }
```

Producer-consumer problem – Understanding

- Semaphore can be more than mutual exclusion!

| | |
|-------|---|
| empty | It represents the number of empty slots. |
| full | It represents the number of occupied slots. |

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6         down(&empty);
7         down(&mutex);
8         insert_item(item);
9         up(&mutex);
10        up(&full);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         down(&full);
6         down(&mutex);
7         item = remove_item();
8         up(&mutex);
9         up(&empty);
10        consume_item(item);
11    }
12 }
```


Producer-consumer problem – question

Question.

Can we swap Lines 6 & 7 of the producer?

Let us simulate what will happen with the modified code!

Shared object

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6*      down(&mutex);
7*      down(&empty);
8         insert_item(item);
9         up(&mutex);
10        up(&full);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         down(&full);
6         down(&mutex);
7         item = remove_item();
8         up(&mutex);
9         up(&empty);
10        consume_item(item);
11    }
12 }
```

Producer-consumer problem – question

Producer

running until Line
10

Consumer

We are showing the value of the semaphores before the producer is suspended.

mutex = 1

empty = 0

full = N

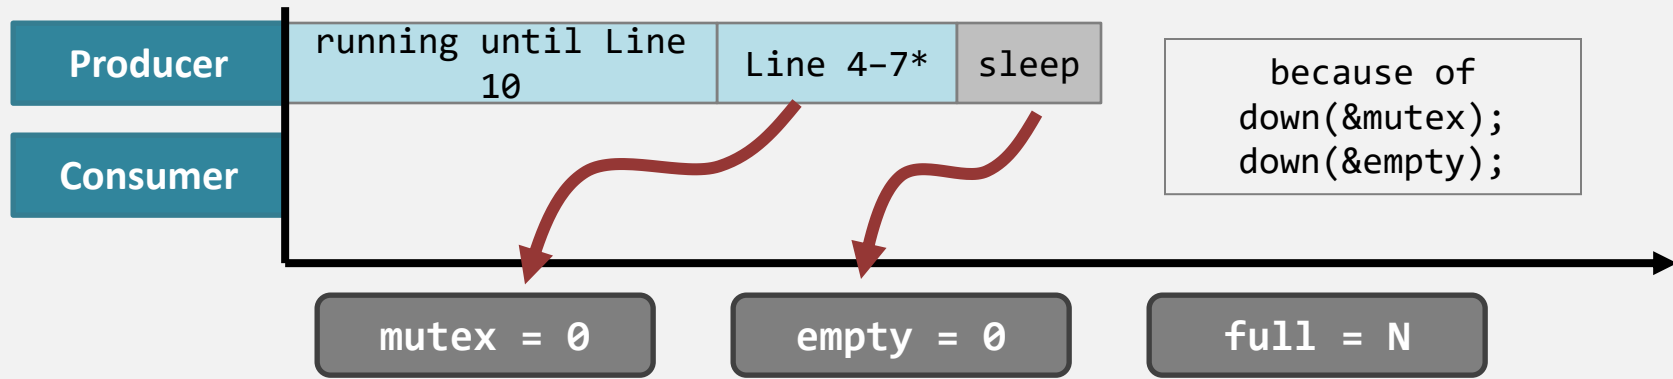
Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6*      down(&mutex);
7*      down(&empty);
8         insert_item(item);
9         up(&mutex);
10        up(&full);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         down(&full);
6         down(&mutex);
7         item = remove_item();
8         up(&mutex);
9         up(&empty);
10        consume_item(item);
11    }
12 }
```

Producer-consumer problem – question



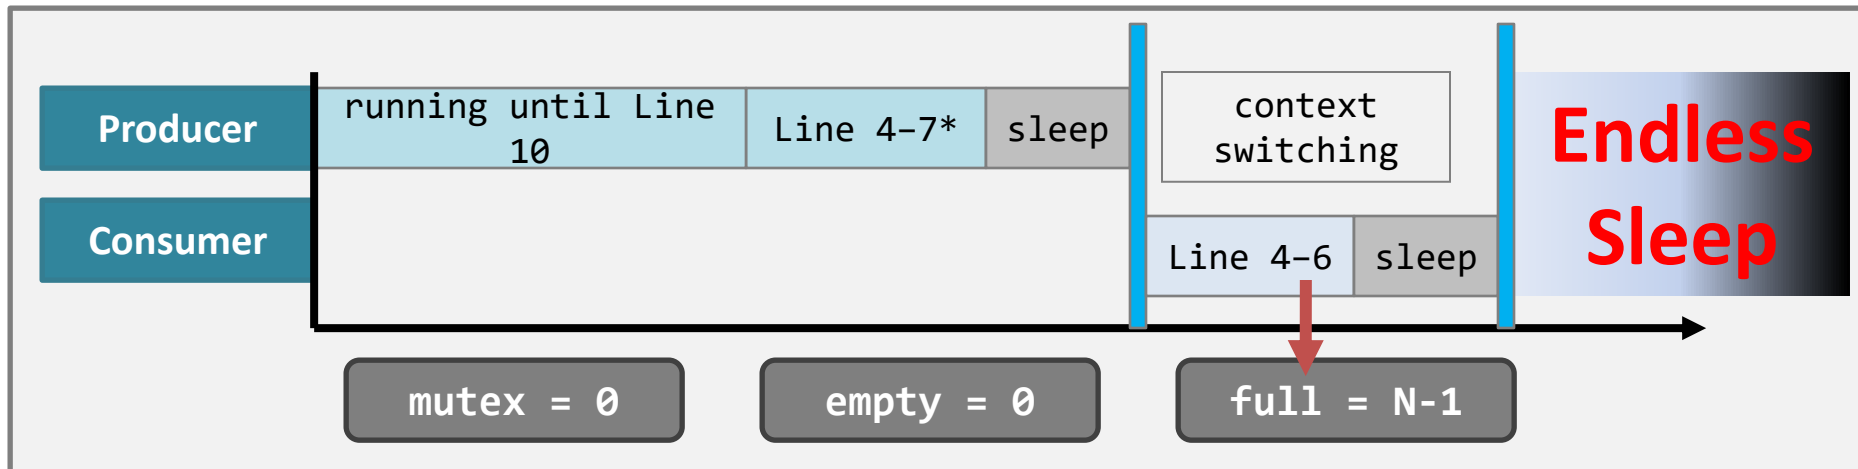
Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6*      down(&mutex);
7*      down(&empty);
8         insert_item(item);
9         up(&mutex);
10        up(&full);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         down(&full);
6         down(&mutex);
7         item = remove_item();
8         up(&mutex);
9         up(&empty);
10        consume_item(item);
11    }
12 }
```

Producer-consumer problem – question



Producer function

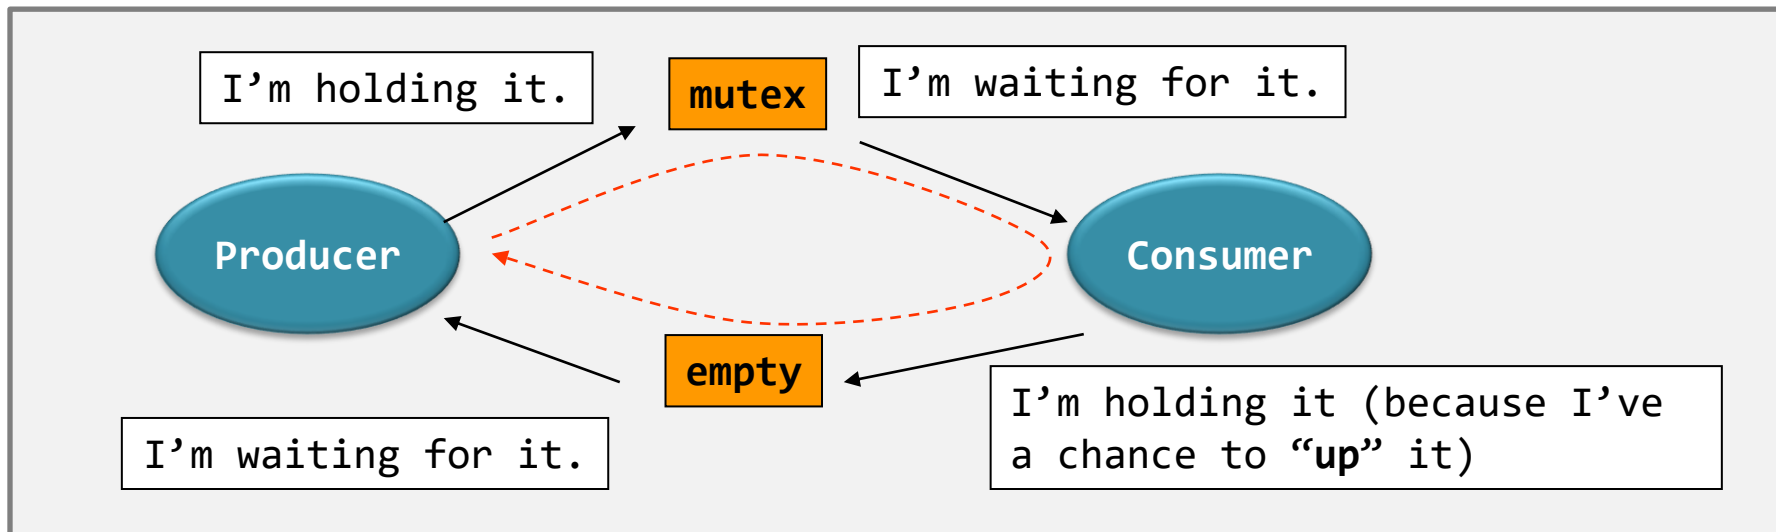
```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6*      down(&mutex);
7*      down(&empty);
8         insert_item(item);
9         up(&mutex);
10        up(&full);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         down(&full);
6         down(&mutex);
7         item = remove_item();
8         up(&mutex);
9         up(&empty);
10        consume_item(item);
11    }
12 }
```

Producer-consumer problem

- **Deadlock** happens when a **circular wait** appears
 - The producer is waiting for the consumer to “up()” the “empty” semaphore, and
 - the consumer is waiting for the producer to “up()” the “mutex” semaphore.




Producer-consumer problem

- **Deadlock** happens when a **circular wait** appears
 - The producer is waiting for the consumer to “**up()**” the “**empty**” semaphore, and
 - the consumer is waiting for the producer to “**up()**” the “**mutex**” semaphore.
- **No progress could be made by all processes + All processes are blocked.**
 - **Implication:** careless implementation of the producer-consumer solution can be disastrous.

Summary on producer-consumer problem

- The problem can be divided into two sub-problems.
 - Mutual exclusion.
 - The buffer is a shared object. Mutual exclusion is needed.
 - Synchronization.
 - Because the buffer's size is bounded, coordination is needed.




Summary on producer-consumer problem

- How to guarantee mutual exclusion?
 - A **binary semaphore** is used as the entry and the exit of the critical sections.
- How to achieve synchronization?
 - Two semaphores are used as **counters** to monitor the status of the buffer.
 - Two semaphores are needed because the two suspension conditions are different.

The Deadlock Problem

Classic IPC problems


- Producer-consumer problem
- **Dining philosopher problem**
- Reader-writer problem




Let's teach them
not to fight.

Dining philosopher – introduction


- 5 philosophers, 5 plates of spaghetti, and 5 chopsticks.
- The jobs of each philosopher are
 - to think and
 - to eat: They **need exactly two chopsticks** in order to eat the spaghetti.
- Question: how to construct a synchronization protocol such that
 - they will not result in any **deadlocking scenarios**, and
 - they will not be **starved to death**




Dining philosopher – introduction




Philosophers



Chopsticks



Spaghetti

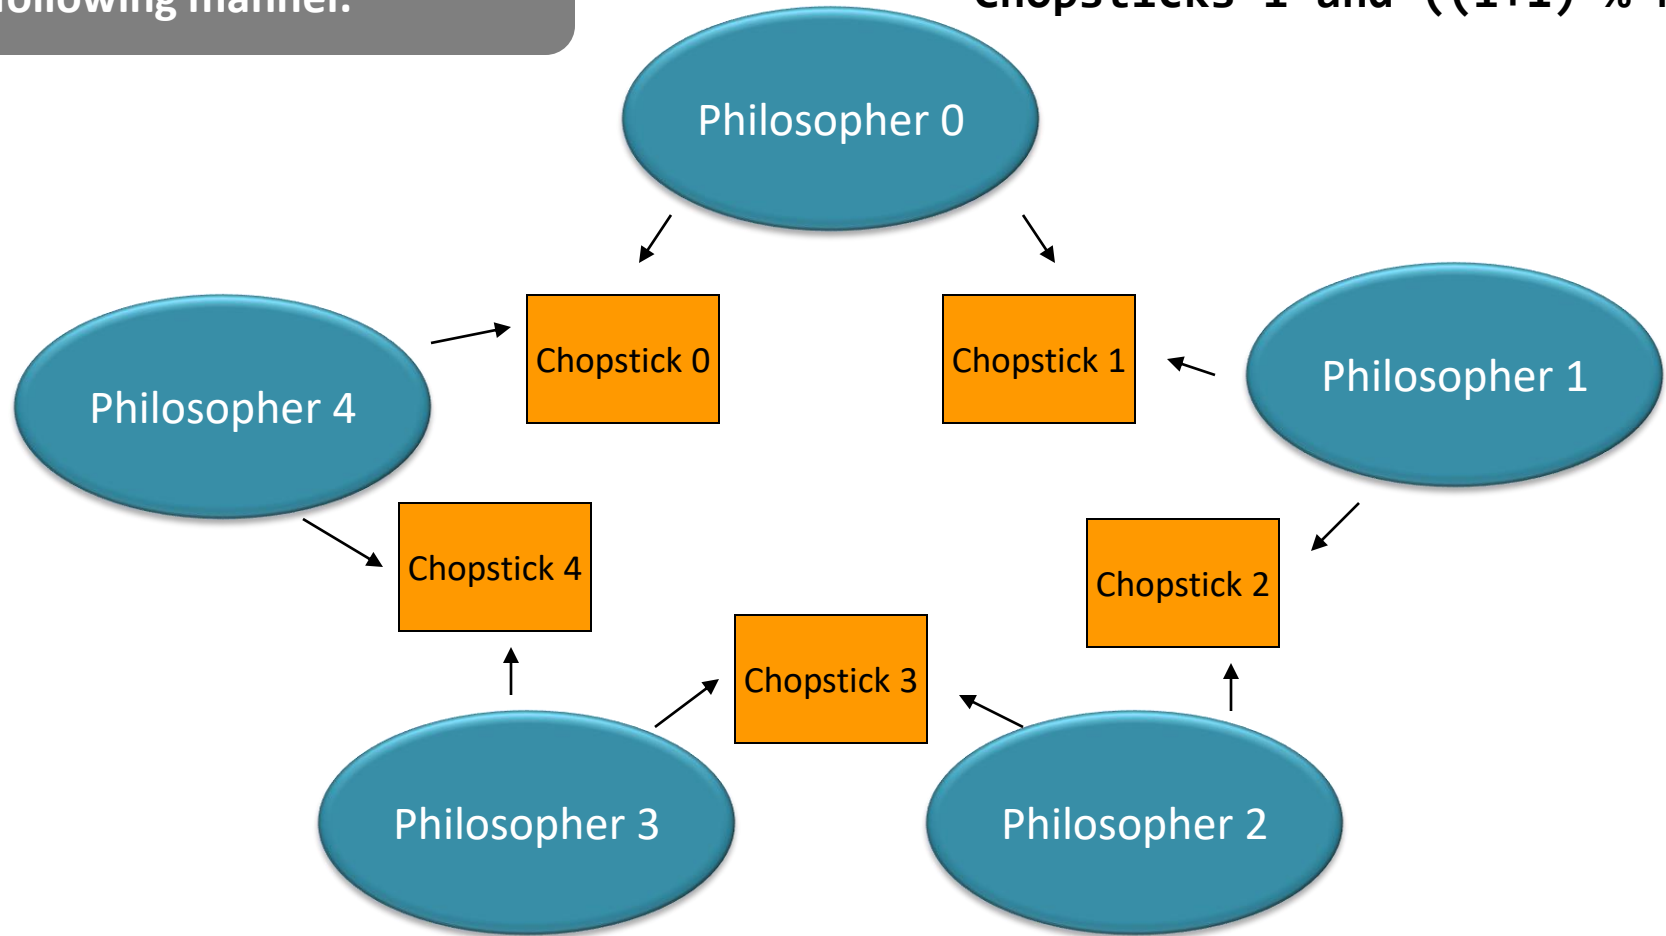


Consider to have
infinite supply.

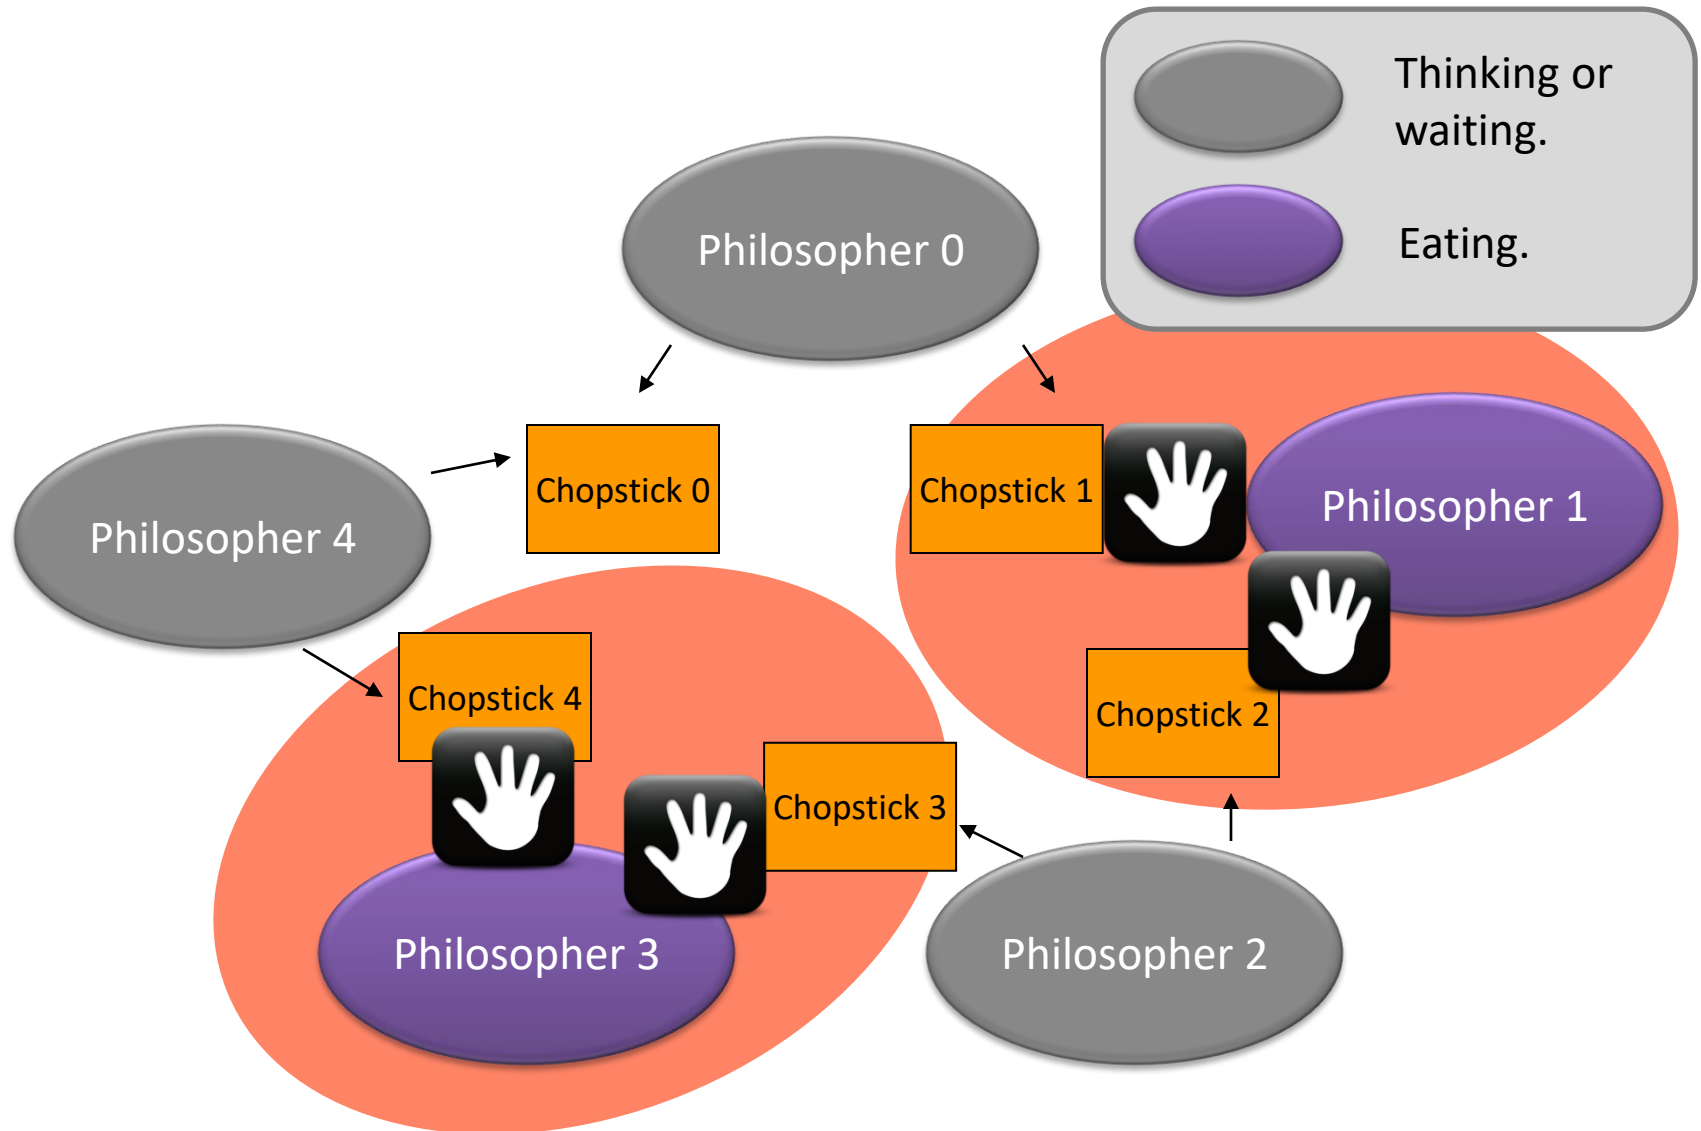
Dining philosopher – introduction

The chopsticks are arranged in the following manner.

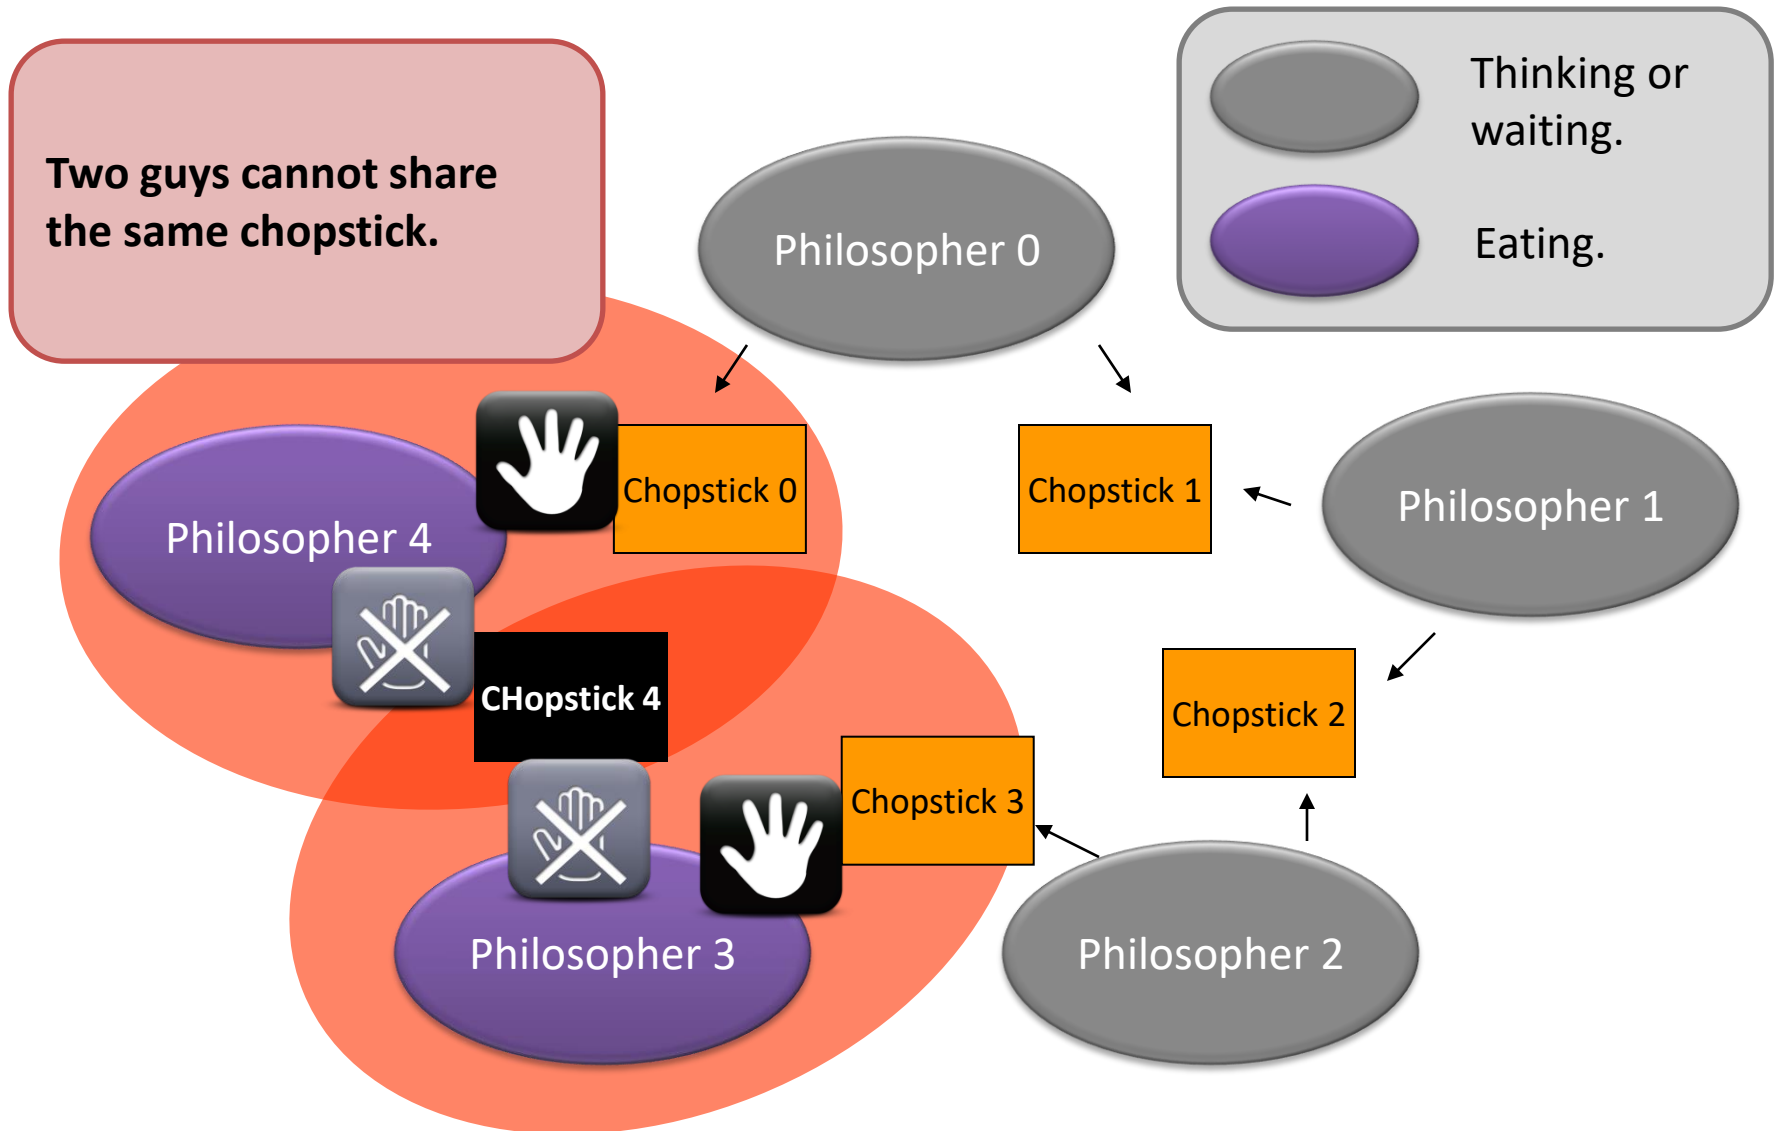
Philosopher i needs
Chopsticks i and $((i+1) \% N)$;



Dining philosopher – introduction



Dining philosopher – introduction



Dining philosopher – requirement #1

- **Mutual exclusion**
 - What if there is no mutual exclusion?
 - Then: while you're eating, the two men besides you will and must **steal all your chopsticks!**
- Let's propose the following solution:
 - When you are hungry, you have to check if anyone is using the chopstick that you need.
 - If yes, you have to wait.
 - If no, **seize both chopsticks.**
 - After eating, put down all your chopsticks.

Dining philosopher – meeting requirement #1?

Shared object

```
#define N 5  
semaphore chop[N];
```

A quick question: what should be initial values?

Helper Functions

```
void take(int i) {  
    down(&chop[i]);  
}
```

```
void put(int i) {  
    up(&chop[i]);  
}
```

Section
Entry

Critical
Section

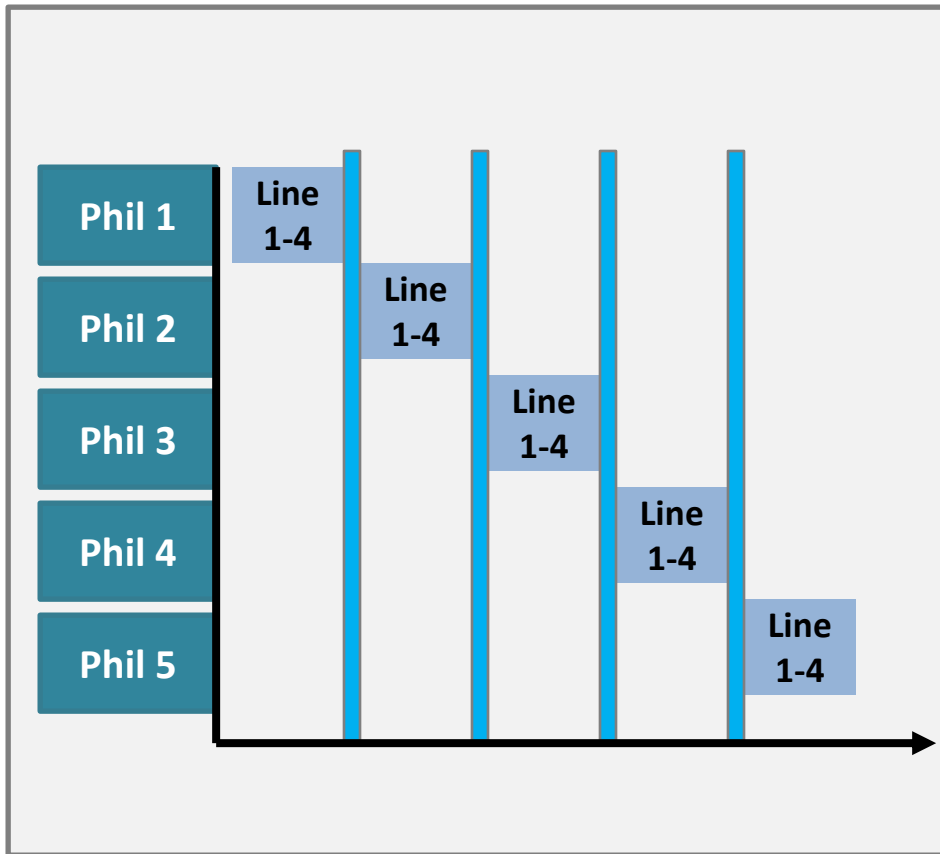
Section
Exit

Main Function

```
1 void philosopher(int i) {  
2     while (TRUE) {  
3         think();  
4         take(i);  
5         take((i+1) % N);  
6         eat();  
7         put(i);  
8         put((i+1) % N);  
9     }  
10 }
```


Dining philosopher – meeting requirement #1?

Final Destination: Deadlock!



Main Function

```
1 void philosopher(int i) {  
2     while (TRUE) {  
3         think();  
4         take(i);  
5         take((i+1) % N);  
6         eat();  
7         put(i);  
8         put((i+1) % N);  
9     }  
10 }
```

Dining philosopher – requirement #2

- **Synchronization**
 - Should avoid any **potential deadlocking execution order**.
- How about the following suggestions:
 - First, a philosopher **takes a chopstick**.
 - If a philosopher finds that he cannot take the second one, then he should **put down the first chopstick**.
 - Then, the philosopher **goes to sleep** for a while.
 - Again, the philosopher tries to get both chopsticks until both ones are seized.

Dining philosopher – meeting requirement #2?

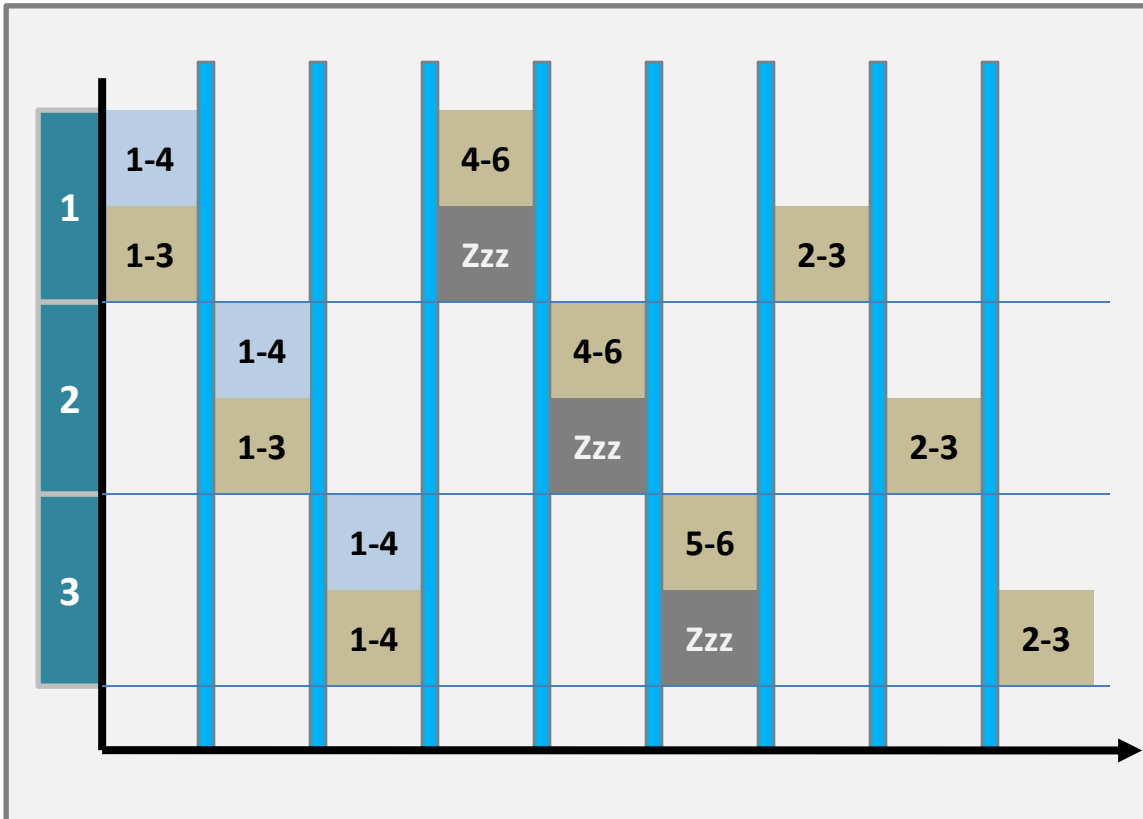
The code: meeting requirement #2?

```
1 void philosopher(int i) {  
2     while (TRUE) {  
3         think();  
4         take(i);  
5         eat();  
6         up(&chop[i]);  
7         up(&chop[(i+1)%N]);  
8     }  
9 }
```

```
1 void take(int i) {  
2     while(TRUE) {  
3         down(&chop[i]);  
4         if (isUsed((i+1)%N)) {  
5             up(&chop[i]);  
6             sleep(1);  
7         }  
8         else {  
9             down(&chop[(i+1)%N]);  
10            break;  
11        }  
12    }  
13 }
```

Dining philosopher – meeting requirement #2?

Potential Problem: Philosophers are all busy
but no progress were made!



Assume N = 3 (because the space is limited)

```
1 void take(int i) {
2   while(TRUE) {
3     down(&chop[i]);
4     if (isUsed((i+1)%N)) {
5       up(&chop[i]);
6       sleep(1);
7     }
8     else {
9       down(&chop[(i+1)%N]);
10      break;
11    }
12  }
13 }
```

```
1 void philosopher(int i) {
2   while (TRUE) {
3     think();
4     take(i);
5     eat();
6     up(&chop[i]);
7     up(&chop[(i+1)%N]);
8   }
9 }
```

Dining philosopher – before the final solution.

- Before we present the final solution, let's see what are the problems that we have.

Problems

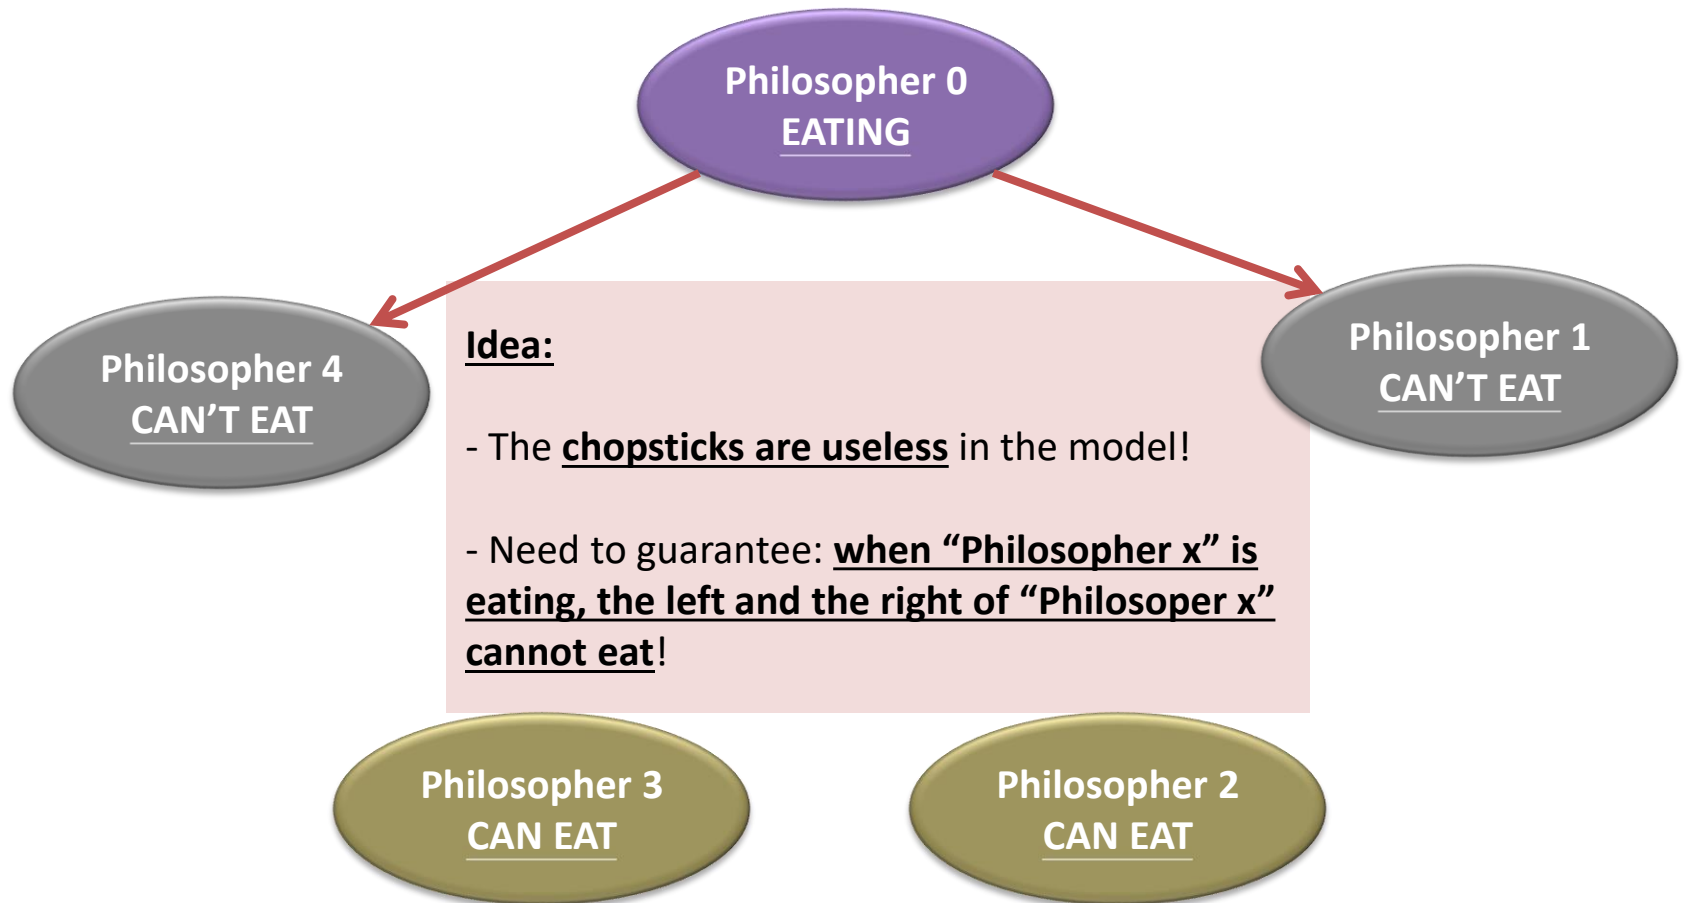
Model a chopstick as a semaphore is intuitive, but is not working.

The problem is that we are afraid to “**down()**”, as that may lead to a deadlock.

Using **sleep()** to avoid deadlock is effective, yet bringing another problem.

We can always create an execution order that keeps all the philosophers busy, but without useful output.

Dining philosopher – before the final solution.



Dining philosopher – the final solution.

Shared object

```
#define N 5
#define LEFT  ((i+N-1) % N)
#define RIGHT  ((i+1) % N)

int state[N];
semaphore mutex = 1;
semaphore s[N];
```

Main function

```
1 void philosopher(int i) {
2     think();
3     take(i);
4     eat();
5     put(i);
6 }
```

Section entry

```
1 void take(int i) {
2     down(&mutex);
3     state[i] = HUNGRY;
4     test(i);
5     up(&mutex);
6     down(&s[i]);
7 }
```

Section exit

```
1 void put(int i) {
2     down(&mutex);
3     state[i] = THINKING;
4     test(LEFT);
5     test(RIGHT);
6     up(&mutex);
7 }
```

I will explain the
code later.

Extremely important helper function

```
1 void test(int i) {
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
3         state[i] = EATING;
4         up(&s[i]);
5     }
6 }
```

Dining philosopher – the final solution.

```
Shared object
#define N 5
#define LEFT  ((i+N-1) % N)
#define RIGHT ((i+1) % N)

int state[N];
semaphore mutex = 1;
semaphore s[N];
```

Going “left” and “right” in a circular manner.

The states of the philosophers, including “**EATING**”, “**THINKING**”, and “**HUNGRY**”.

Remember, this is shared array.

To guarantee mutual exclusive access to the “**state[N]**” array.

Guess:

What is the meaning of the semaphore **s[N]**?

To fulfill the synchronization requirement.

Question. What are the initial values of the “**s[N]**” array?

Dining philosopher – the final solution.

Shared object

```
#define N 5
#define LEFT  ((i+N-1) % N)
#define RIGHT ((i+1) % N)

int state[N];
semaphore mutex = 1;
semaphore s[N];
```

Section entry

```
1 void take(int i) {
2     down(&mutex);
3     state[i] = HUNGRY;
4     test(i);
5     up(&mutex);
6     down(&s[i]);
7 }
```

Question. What are they doing?

If both chopsticks are available,
I eat. Else, I sleep.

Extremely important helper function

```
1 void test(int i) {
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
3         state[i] = EATING;
4         up(&s[i]);
5     }
6 }
```

If they are eating, I can't be eating.

Dining philosopher – the final solution.

Try to let the one on the **left of the caller** to eat.

Try to let the one on the **right of the caller** to eat.

Section exit

```
1 void put(int i) {  
2     down(&mutex);  
3     state[i] = THINKING;  
4     test(LEFT);  
5     test(RIGHT);  
6     up(&mutex);  
7 }
```

Extremely important helper function

```
1 void test(int i) {  
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {  
3         state[i] = EATING;  
4         up(&s[i]);  
5     }  
6 }
```

Wake up the one who can eat!

Dining philosopher – the final solution.

An illustration: How can
Philosopher 1 start eating?

Philosopher 0
THINKING

Philosopher 4
THINKING

Note: no chopsticks objects
will be shown in this
illustration because we
don't need them now.

Philosopher 1
THINKING

Philosopher 3
THINKING

Philosopher 2
THINKING

Dining philosopher – the final solution.

Section entry

```
1 void take(int i) {  
2     down(&mutex);  
3     state[i] = HUNGRY;  
4     test(i);  
5     up(&mutex);  
6     down(&s[i]);  
7 }
```

Call take();

Philosopher 0
HUNGRY

To LEFT:
are you "EATING"?

To RIGHT:
are you "EATING"?

Philosopher 4
THINKING

Philosopher 1
THINKING

Philosopher 3
THINKING

Philosopher 2
THINKING

Dining philosopher – the final solution.

Section entry

```
1 void take(int i) {  
2     down(&mutex);  
3     state[i] = HUNGRY;  
4     test(i);  
5     up(&mutex);  
6     down(&s[i]);  
7 }
```

Call `take()`;

Philosopher 0
HUNGRY

To LEFT:
are you “EATING”?

Philosopher 4
THINKING

To RIGHT:
are you “EATING”?

Philosopher 1
THINKING

Calling `take()`.
but, it is blocked.
Why?

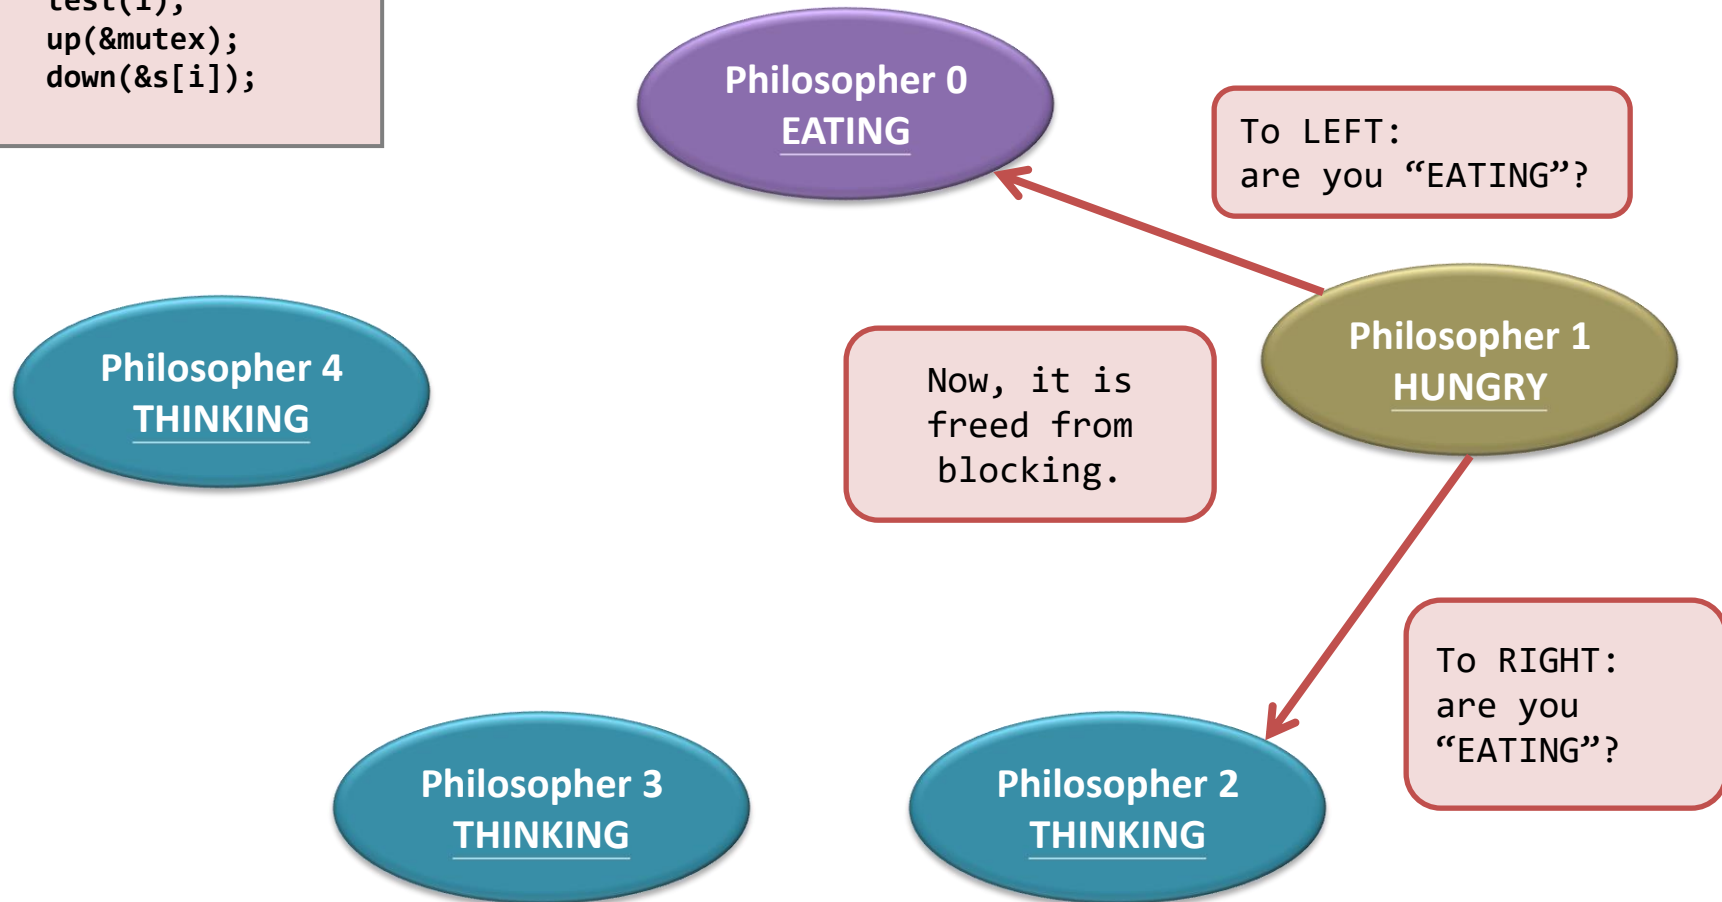
Philosopher 3
THINKING

Philosopher 2
THINKING

Dining philosopher – the final solution.

Section entry

```
1 void take(int i) {  
2     down(&mutex);  
3     state[i] = HUNGRY;  
4     test(i);  
5     up(&mutex);  
6     down(&s[i]);  
7 }
```



Dining philosopher – the final solution.

Section entry

```
1 void take(int i) {  
2     down(&mutex);  
3     state[i] = HUNGRY;  
4     test(i);  
5     up(&mutex);  
6     down(&s[i]);  
7 }
```

Philosopher 0
EATING

Philosopher 4
THINKING

Philosopher 1
HUNGRY

To RIGHT:
are you
“EATING”?

To LEFT:
are you
“EATING”?

Blocked;
because of
`down(&s[1]);`

Philosopher 3
HUNGRY

Philosopher 2
THINKING

Dining philosopher – the final solution.

Section entry

```
1 void take(int i) {  
2     down(&mutex);  
3     state[i] = HUNGRY;  
4     test(i);  
5     up(&mutex);  
6     down(&s[i]);  
7 }
```

Philosopher 0
EATING

Philosopher 4
THINKING

Philosopher 1
HUNGRY

Blocked;
because of
down(&s[1]);

Philosopher 3
EATING

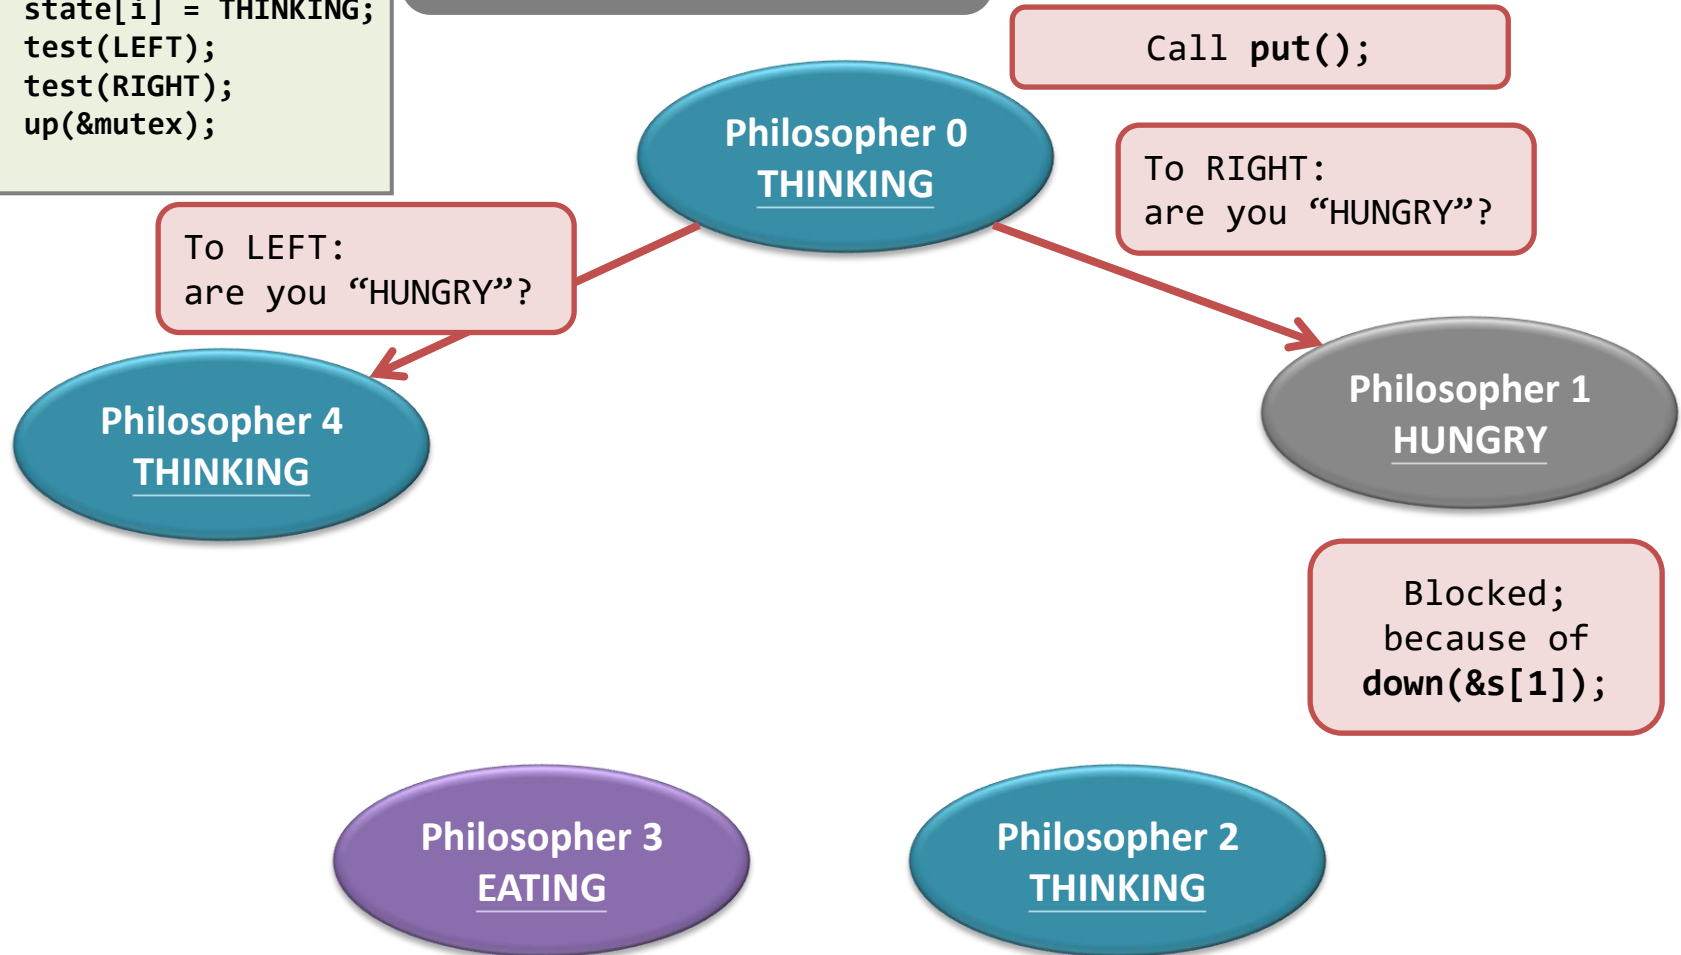
Philosopher 2
THINKING

Dining philosopher – the final solution.

Section exit

```
1 void put(int i) {  
2     down(&mutex);  
3     state[i] = THINKING;  
4     test(LEFT);  
5     test(RIGHT);  
6     up(&mutex);  
7 }
```

An illustration: How can
Philosopher 1 start eating?

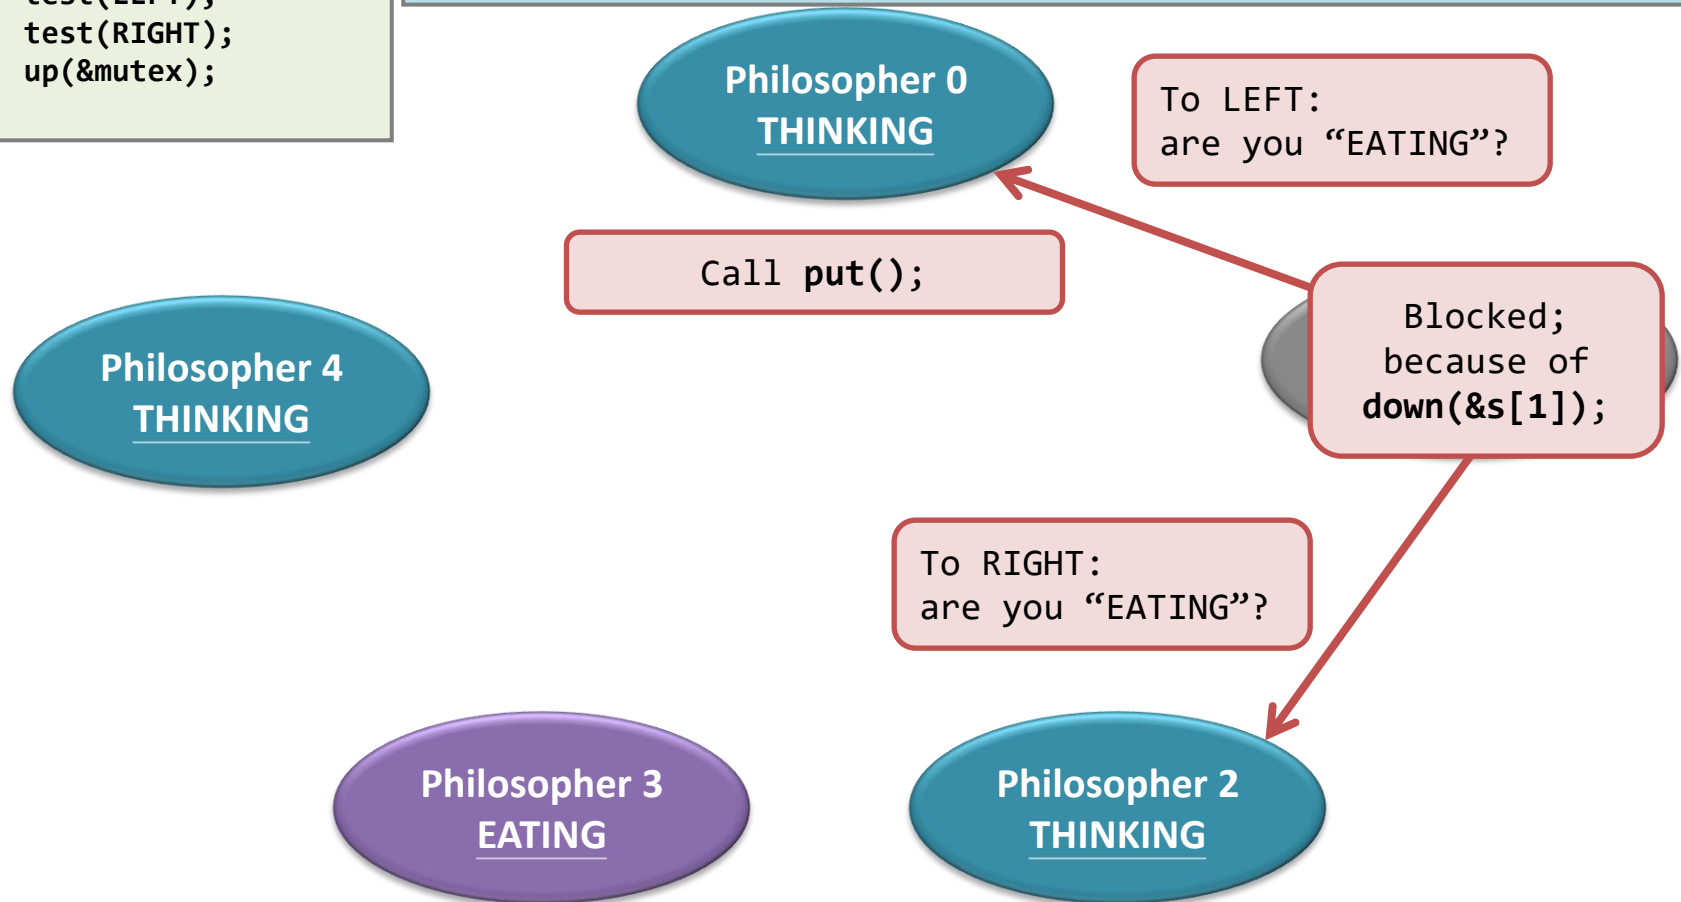


Dining philosopher – the final solution.

Section exit

```
1 void put(int i) {  
2     down(&mutex);  
3     state[i] = THINKING;  
4     test(LEFT);  
5     test(RIGHT);  
6     up(&mutex);  
7 }
```

```
1 void test(int i) {  
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {  
3         state[i] = EATING;  
4         up(&s[i]);  
5     }  
6 }
```

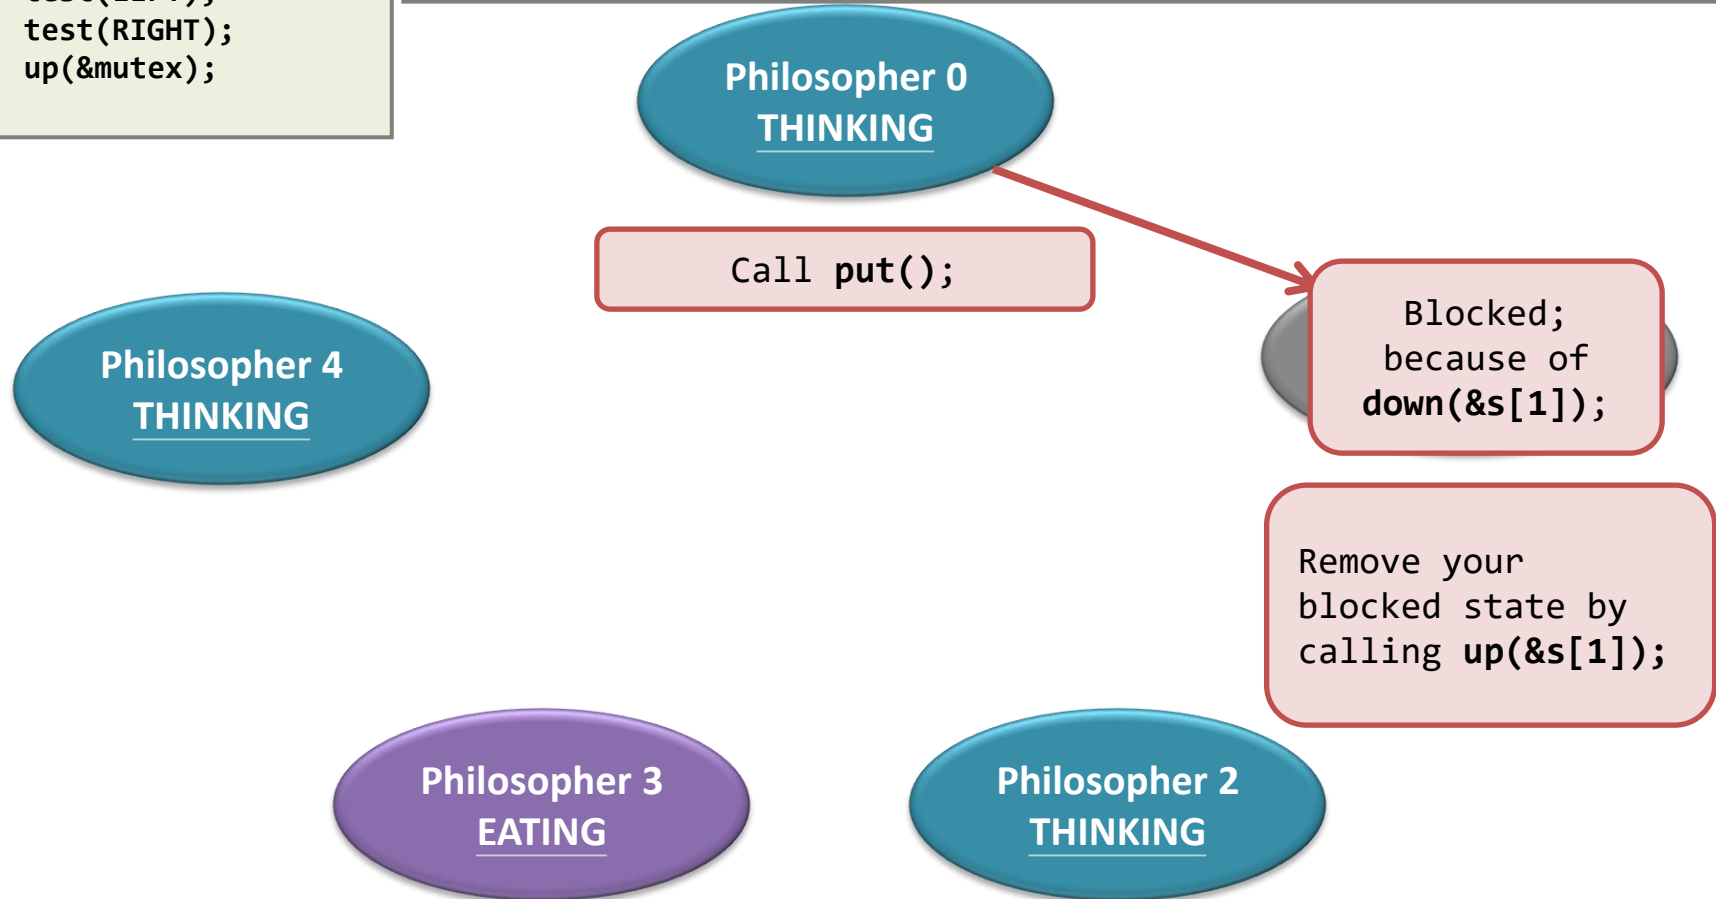


Dining philosopher – the final solution.

Section exit

```
1 void put(int i) {  
2     down(&mutex);  
3     state[i] = THINKING;  
4     test(LEFT);  
5     test(RIGHT);  
6     up(&mutex);  
7 }
```


```
1 void test(int i) {  
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {  
3         state[i] = EATING;  
4         up(&s[i]);  
5     }  
6 }
```



Dining philosopher – the final solution.

Section entry

```
1 void take(int i) {  
2     down(&mutex);  
3     state[i] = HUNGRY;  
4     test(i);  
5     up(&mutex);  
6     down(&s[i]);  
7 }
```



Philosopher 0
THINKING

Philosopher 4
THINKING

Philosopher 1
EATING

Eventually...

Philosopher 3
EATING

Philosopher 2
THINKING

Handwritten Chinese text: s[i] 在 sharing 2 个 筷子 (s[i] is sharing 2 chopsticks)

Dining philosopher - summary

- What is the shared object in the final solution?
 - How to guarantee the mutual exclusion

| Section entry | Section exit |
|--|--|
| <pre>1 void take(int i) { 2 down(&mutex); 3 state[i] = HUNGRY; 4 test(i); 5 up(&mutex); 6 down(&s[i]); 7 }</pre> | <pre>1 void put(int i) { 2 down(&mutex); 3 state[i] = THINKING; 4 test(LEFT); 5 test(RIGHT); 6 up(&mutex); 7 }</pre> |

Dining philosopher - summary

- Think:
 - Why the semaphore $s[N]$ is needed
 - How to set its initial value

Section entry

```
1 void take(int i) {  
2     down(&mutex);  
3     state[i] = HUNGRY;  
4     test(i);  
5     up(&mutex);  
6     down(&s[i]);  
7 }
```

Extremely important helper function

```
1 void test(int i) {  
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {  
3         state[i] = EATING;  
4         up(&s[i]);  
5     }  
6 }
```

Dining philosopher - summary


- Solution to IPC problem can be difficult to comprehend.
 - Usually, intuitive methods failed.
 - Depending on time, e.g., `sleep(1)`, does not guarantee a useful solution.
- As a matter of fact, dining philosopher **is not restricted to 5 philosophers.**

The Deadlock Problem

Classic IPC problems


- Producer-consumer problem
- Dining philosopher problem
- Reader-writer problem

Let's teach them
not to fight.



Reader-writer problem – introduction

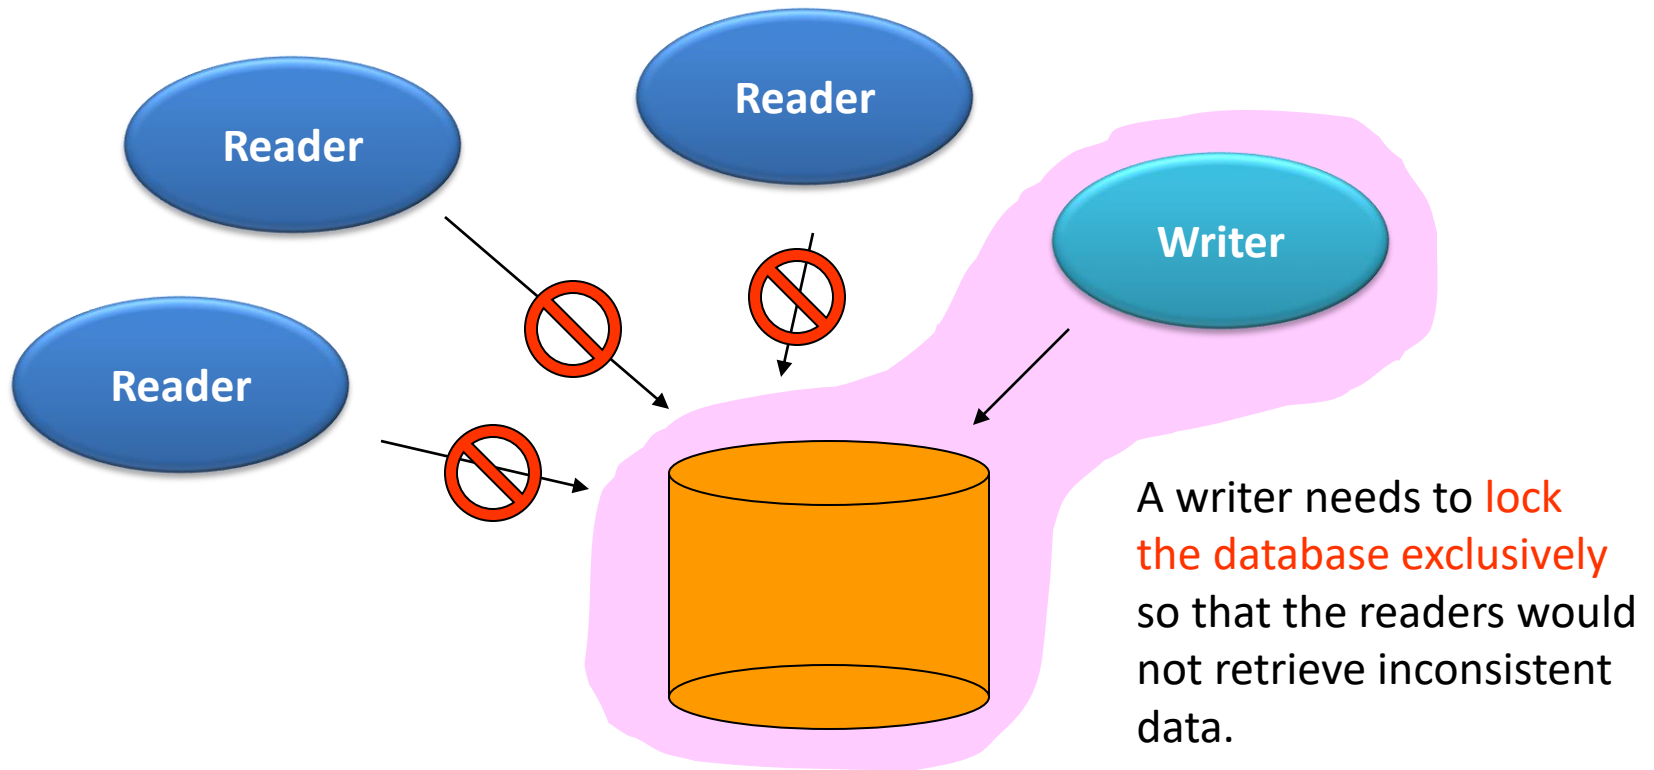
- It is a concurrent database problem.



Readers are allowed to read the content of the database concurrently.

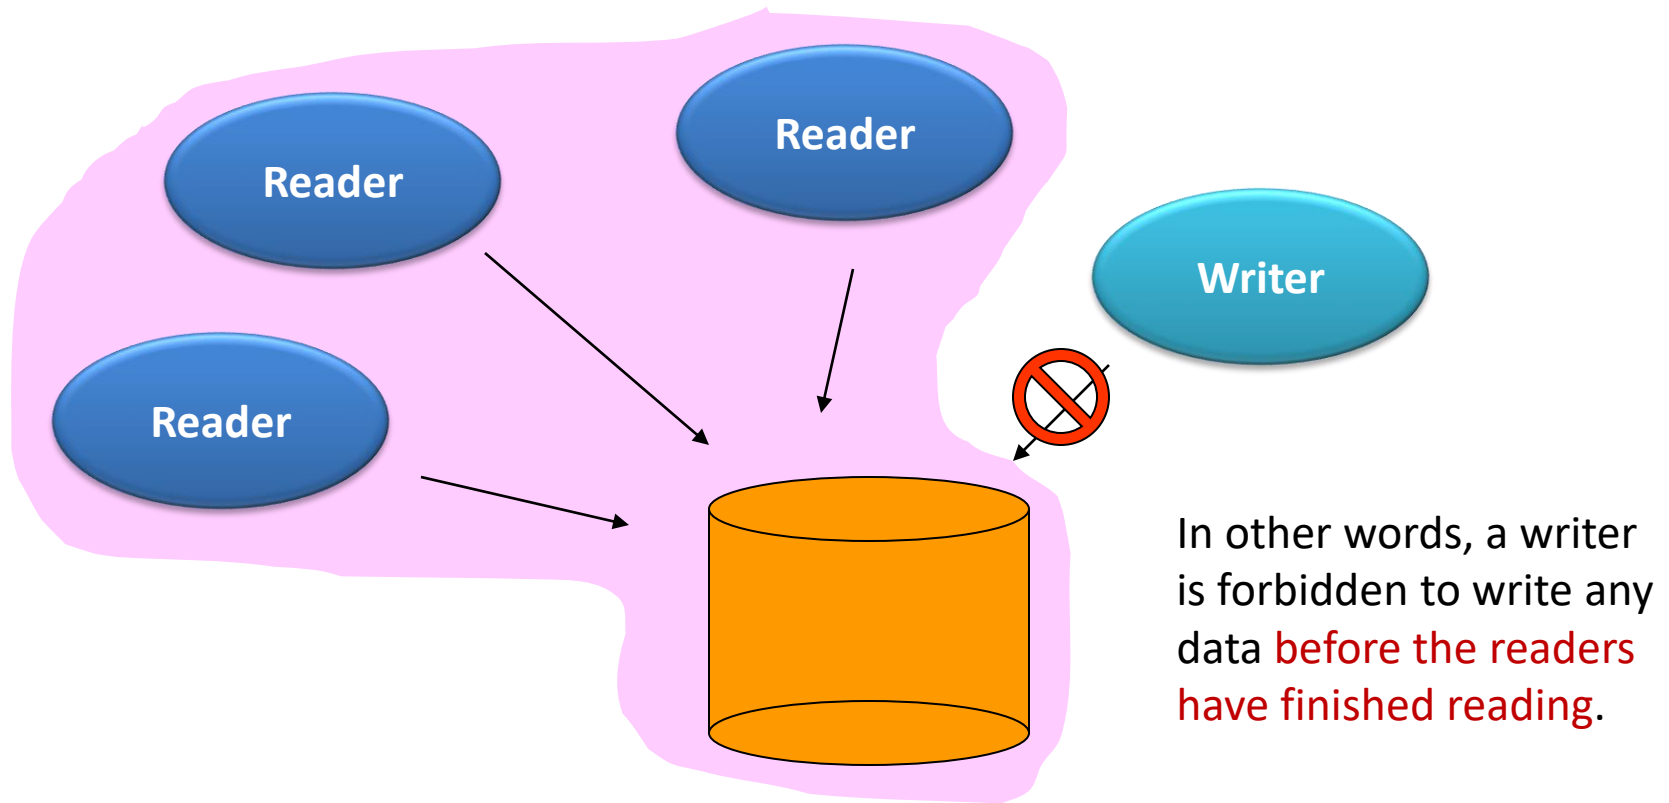
Reader-writer problem – introduction

- It is a concurrent database problem.



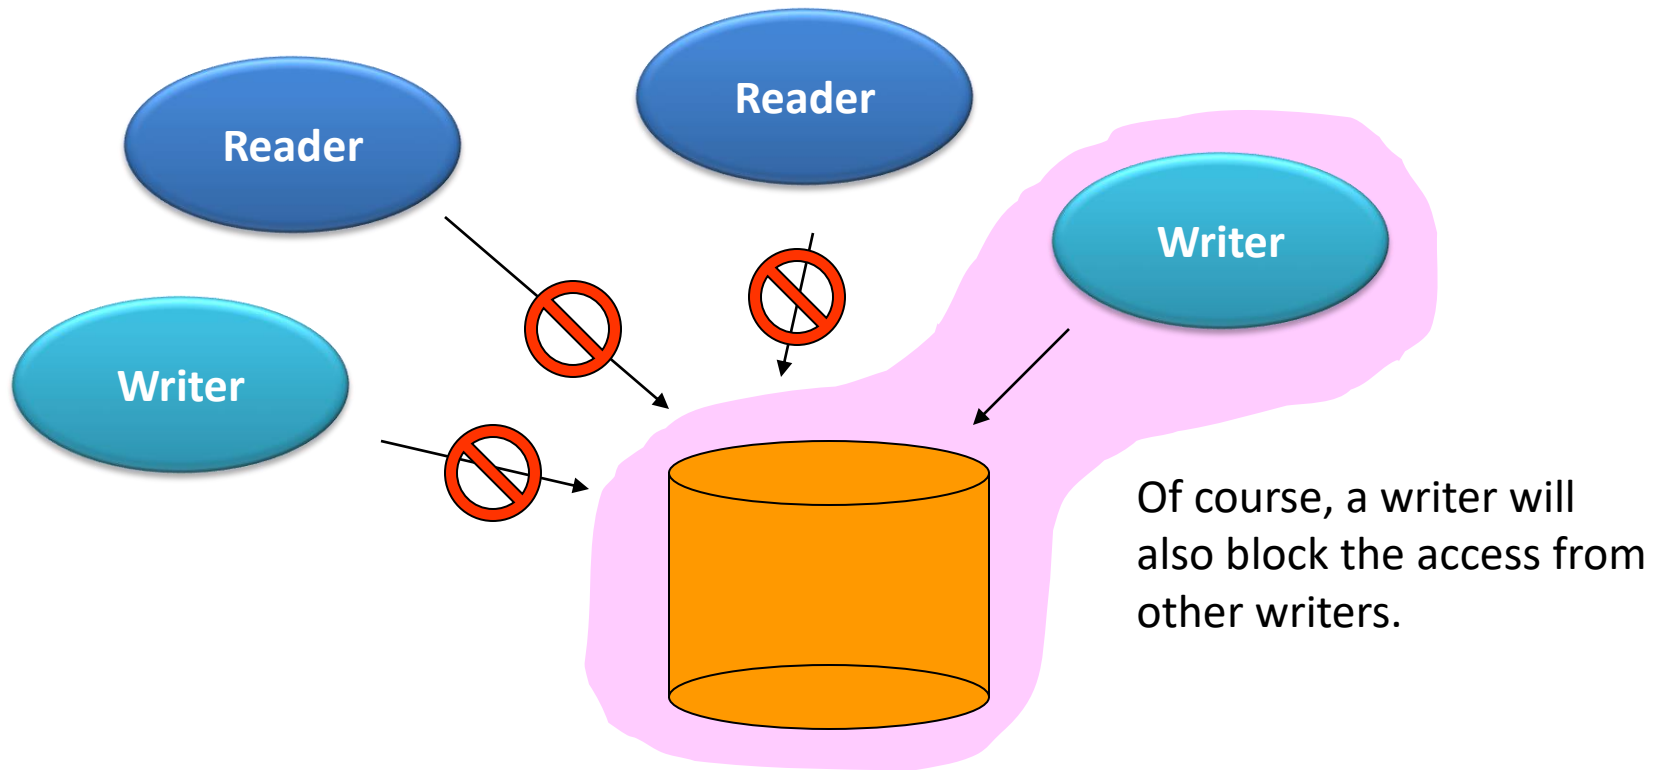
Reader-writer problem – introduction

- It is a concurrent database problem.



Reader-writer problem – introduction

- It is a concurrent database problem.

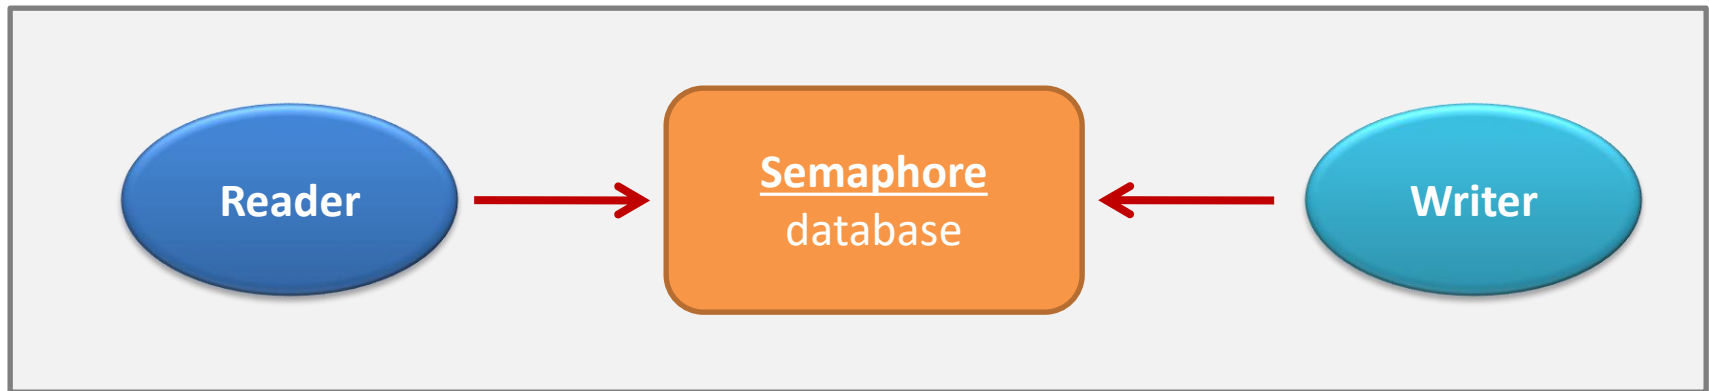


Reader-writer problem – subproblems

- A mutual exclusion problem.
 - The database is a shared object.
- A synchronization problem.
 - **Rule 1.** While a reader is reading, other readers is allowed to read the database.
 - **Rule 2.** While a reader is reading, no writers is allowed to write to the database.
 - **Rule 3.** While a writer is writing, no writers and readers are allowed to access the database.
- A concurrency problem.
 - **Simultaneous access for multiple readers** is allowed and must be guaranteed.

Reader-writer problem – solution outline

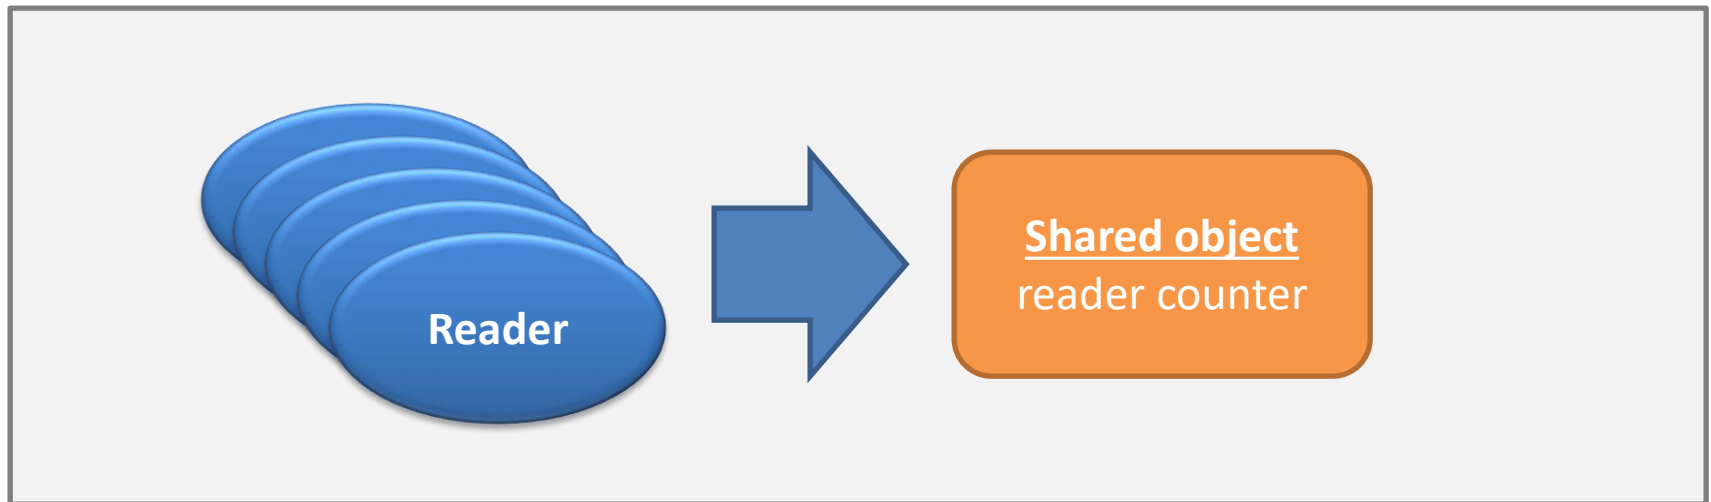
- **Mutual exclusion**: relate the readers and the writers to one semaphore.
 - This guarantees **no readers and writers** could proceed to their critical sections at the same time.
 - This also guarantees **no two writers** could proceed to their critical sections at the same time.



Reader-writer problem – solution outline

- Readers' concurrency

- The **first reader coming** to the system “**down()**” the “**database**” semaphore.
- The **last reader leaving** the system “**up()**” the “**database**” semaphore.



Reader-writer problem – final solution

Shared object

```
semaphore db    = 1;
semaphore mutex = 1;
int read_count  = 0;
```

Writer function

```
1 void writer(void) {
2   while(TRUE) {
```

Section Entry

```
   prepare_write();
   down(&db);
```

Critical Section

```
   write_database();
```

Section Exit

```
   up(&db);
```

```
7   }
8 }
```

Reader Function

```
1 void reader(void) {
2   while(TRUE) {
```

Section Entry

```
   down(&mutex);
   read_count++;
5   if(read_count == 1)
6       down(&db);
7   up(&mutex);
```

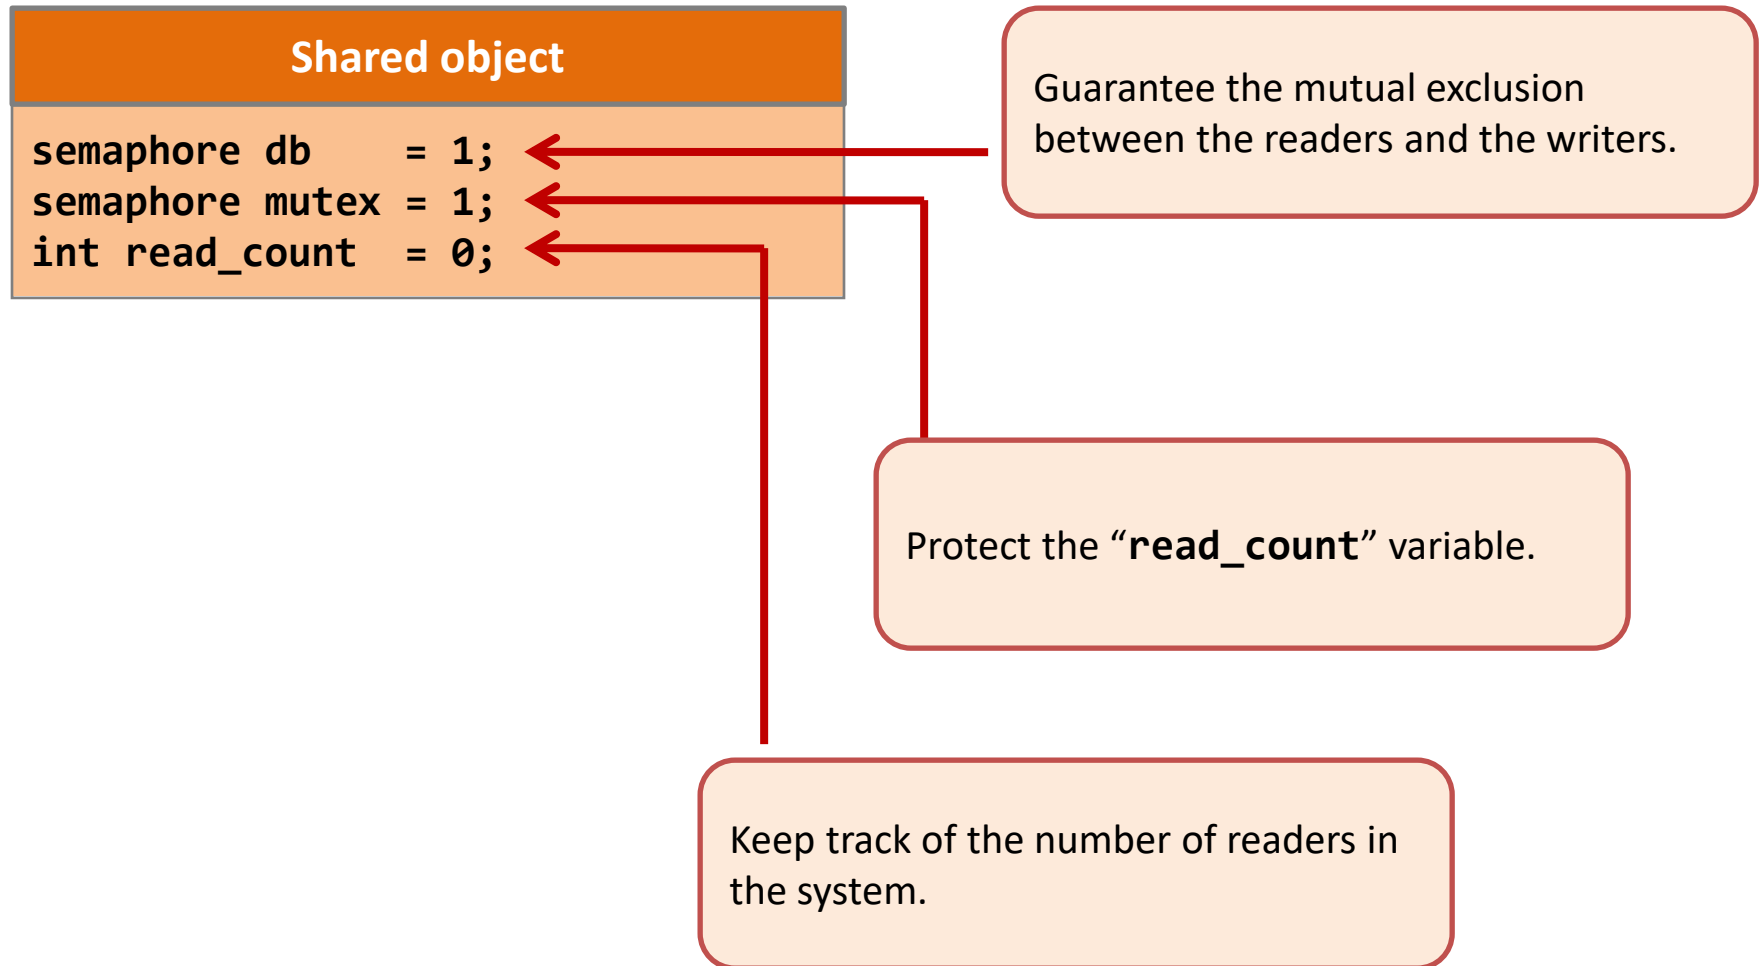
Critical Section

```
   read_database();
```

Section Exit

```
   down(&mutex);
   read_count--;
11  if(read_count == 0)
12      up(&db);
13  up(&mutex);
14  process_data();
15  }
16 }
```


Reader-writer problem – final solution



Reader-writer problem – final solution

Shared object

```
semaphore db    = 1;  
semaphore mutex = 1;  
int read_count  = 0;
```

Writer function

```
1 void writer(void) {  
2   while(TRUE) {  
   Section Entry  prepare_write();  
                  down(&db);  
   Critical Section write_database();  
   Section Exit   up(&db);  
7   }  
8 }
```

The writer is allowed to enter its critical section when no other process is in its critical section (protected by the “**db**” semaphore)

Reader-writer problem – final solution

Shared object

```
semaphore db    = 1;  
semaphore mutex = 1;  
int read_count  = 0;
```

The first reader “**down()**” the “**db**” semaphore so that no writers would be allowed to enter their critical sections.

The last reader “**up()**” the “**db**” semaphore so as to let the writers to enter their critical section.

Reader Function

```
1 void reader(void) {  
2     while(TRUE) {  
3         down(&mutex);  
4         read_count++;  
5         if(read_count == 1)  
6             down(&db);  
7         up(&mutex);  
  
8         read_database();  
  
9         down(&mutex);  
10        read_count--;  
11        if(read_count == 0)  
12            up(&db);  
13        up(&mutex);  
14        process_data();  
15    }  
16 }
```

Reader-writer problem – summary

- This solution does not limit the number of readers and the writers admitted to the system.
 - A realistic database needs this property.
- This solution gives readers a higher priority over the writers.
 - Whenever there are readers, writers must be blocked, not the other way round.
- **What if a writer should be given a higher priority?**

Summary on IPC problems

- The problems have the following properties in common:
 - Multiple processes;
 - Shared and limited resources;
 - Processes have to be synchronized in order to generate useful output;
- The synchronization algorithms have the following requirements in common:
 - Guarantee mutual exclusion;
 - Uphold the correct synchronization among processes;
 - Deadlock-free.

Summary on Ch5

